

Computer Vision

What is computer vision

- Computer vision is a field that includes methods for acquiring, processing, analysing and understanding images
- Duplicate the abilities of human vision by electronically perceiving and understanding an image.
- Image data can take many forms, such as a video sequence, depth images, views from multiple cameras, medical scanner, satellite images etc.

What a computer sees?

COMPUTER VISION VS HUMAN VISION



What we see

What a computer sees?

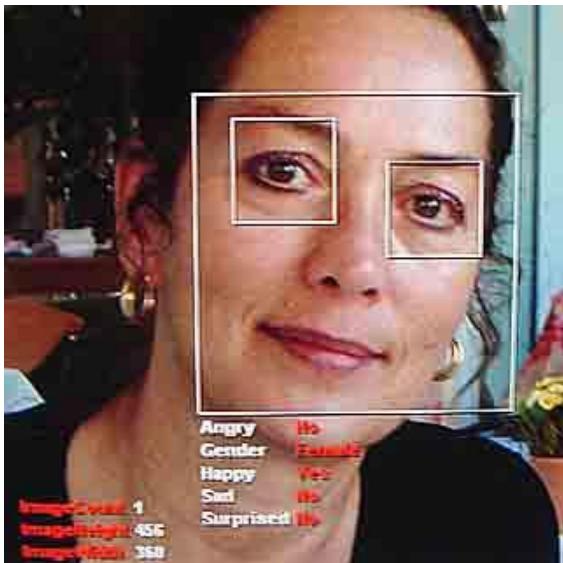
COMPUTER VISION VS HUMAN VISION



What we see

What a computer sees

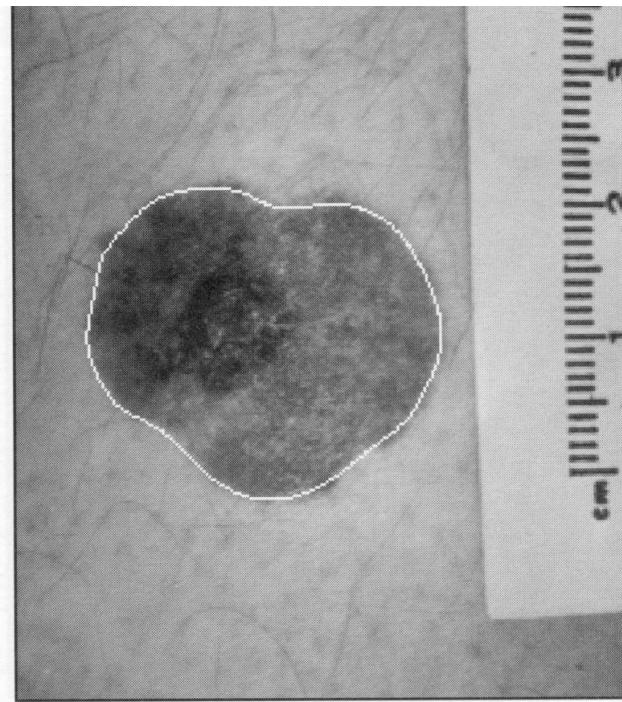
- Learn to recognize faces and expressions.



ALGORITHMS
Support Vector Machines
Adaboost

Medical Applications

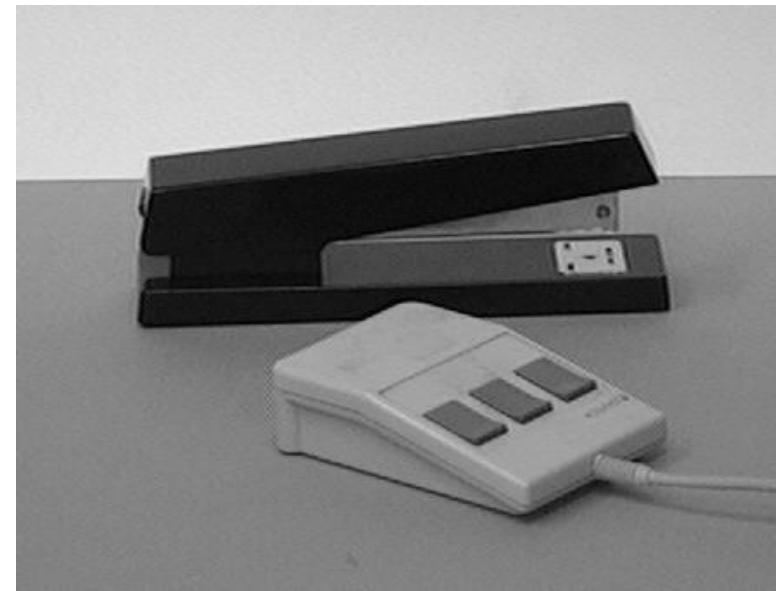
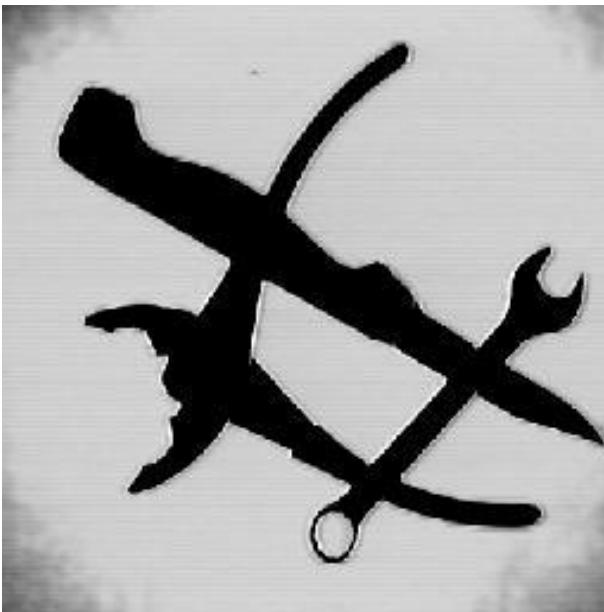
- skin cancer



- breast cancer



Object Recognition

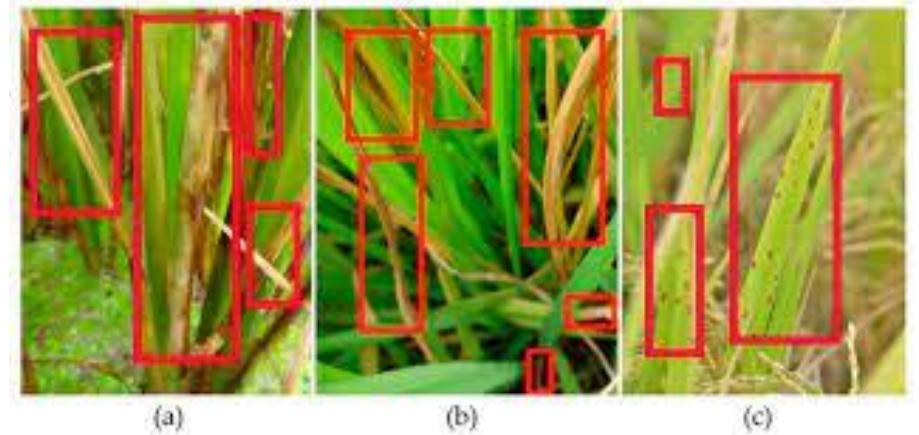


Object Detection



Agriculture

- Crop monitoring, flower detection
plant disease detection
Insect detection



Autonomous Vehicles

- Land, Underwater, Space



Interpretation of Aerial Photography



Interpretation of satellite Imagery



ALGORITHMS
Mask RCNN
Single shot detection
yolo

Mining

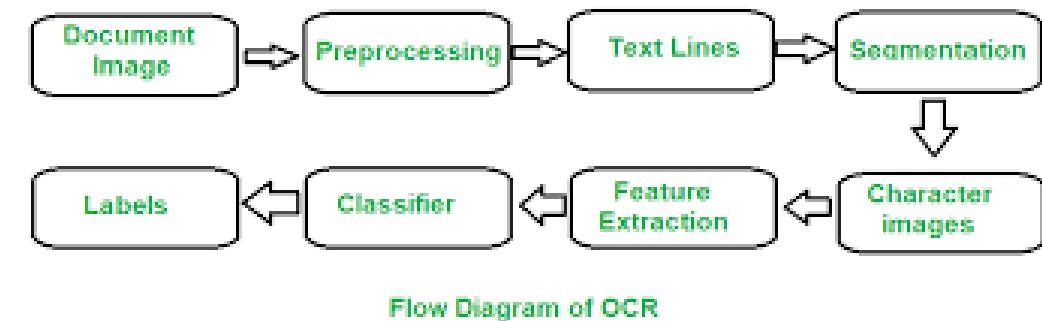


Computer vision and IOT

Character Recognition



pytesseract



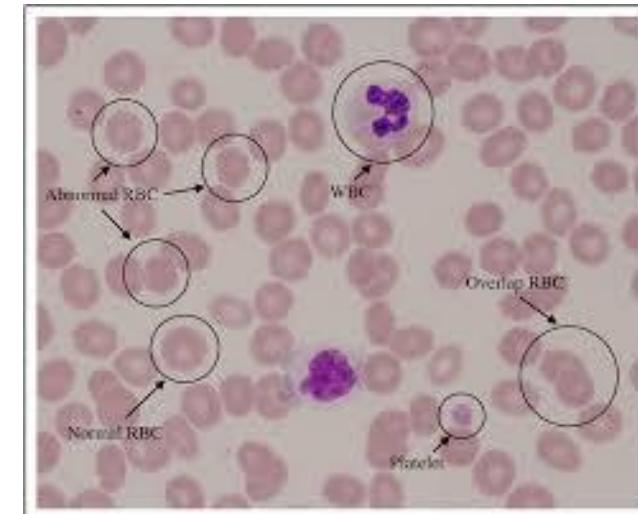
Flow Diagram of OCR

Many Applications: Invoice, id cards, Number plate detection etc

Bio Informatics



Cell classification



Signature recognition

- Recognize signatures by structural similarities which are difficult to quantify.
- does a signature belongs to a specific person, say Tony Blair, or not.



John Brem



Suzanne Pleshette

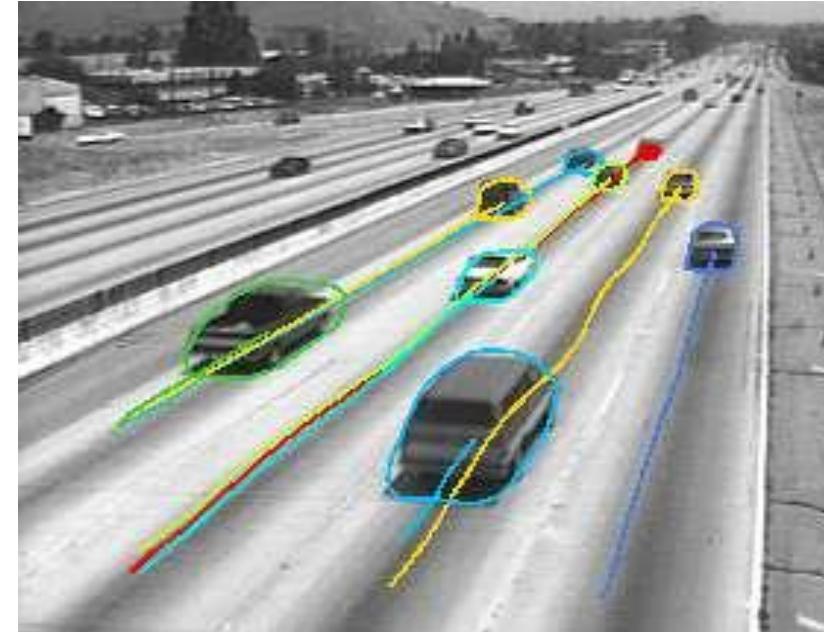
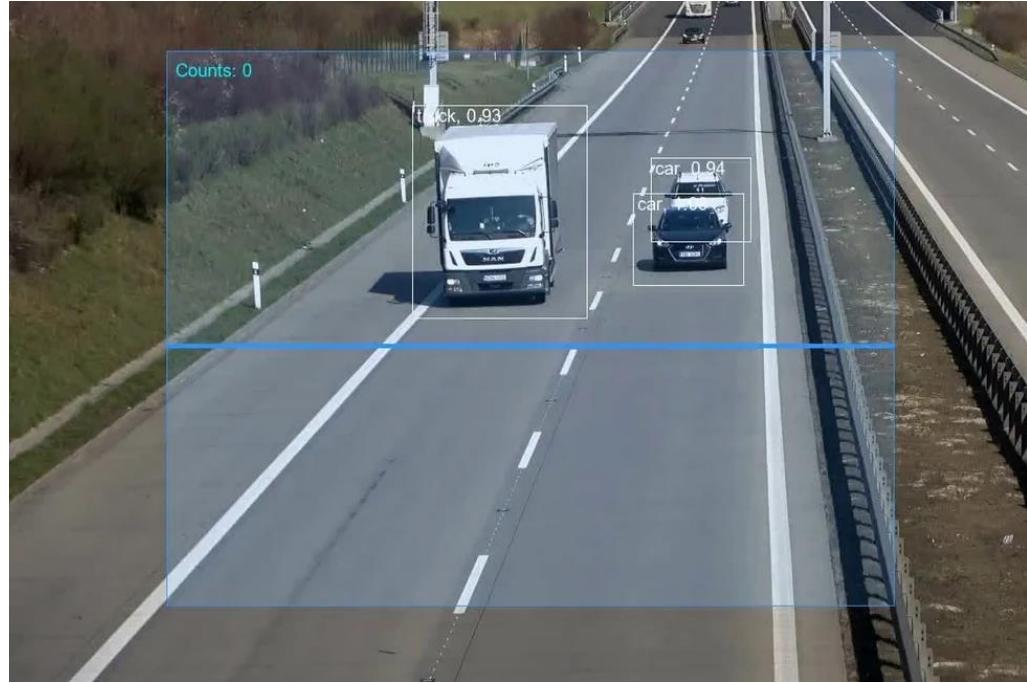


K. James



Glynis Johns

Traffic Monitoring



Customer tracking in Malls



Strategically placed counting devices throughout a retail store can gather data through machine learning processes about where customers spend their time, and for how long.

Customer analytics can improve retail stores' understanding of consumer interaction and improve store layout optimization.

Image transformation using GANS



Create faces that never exists

Art using GANS



Defect detection in Industry

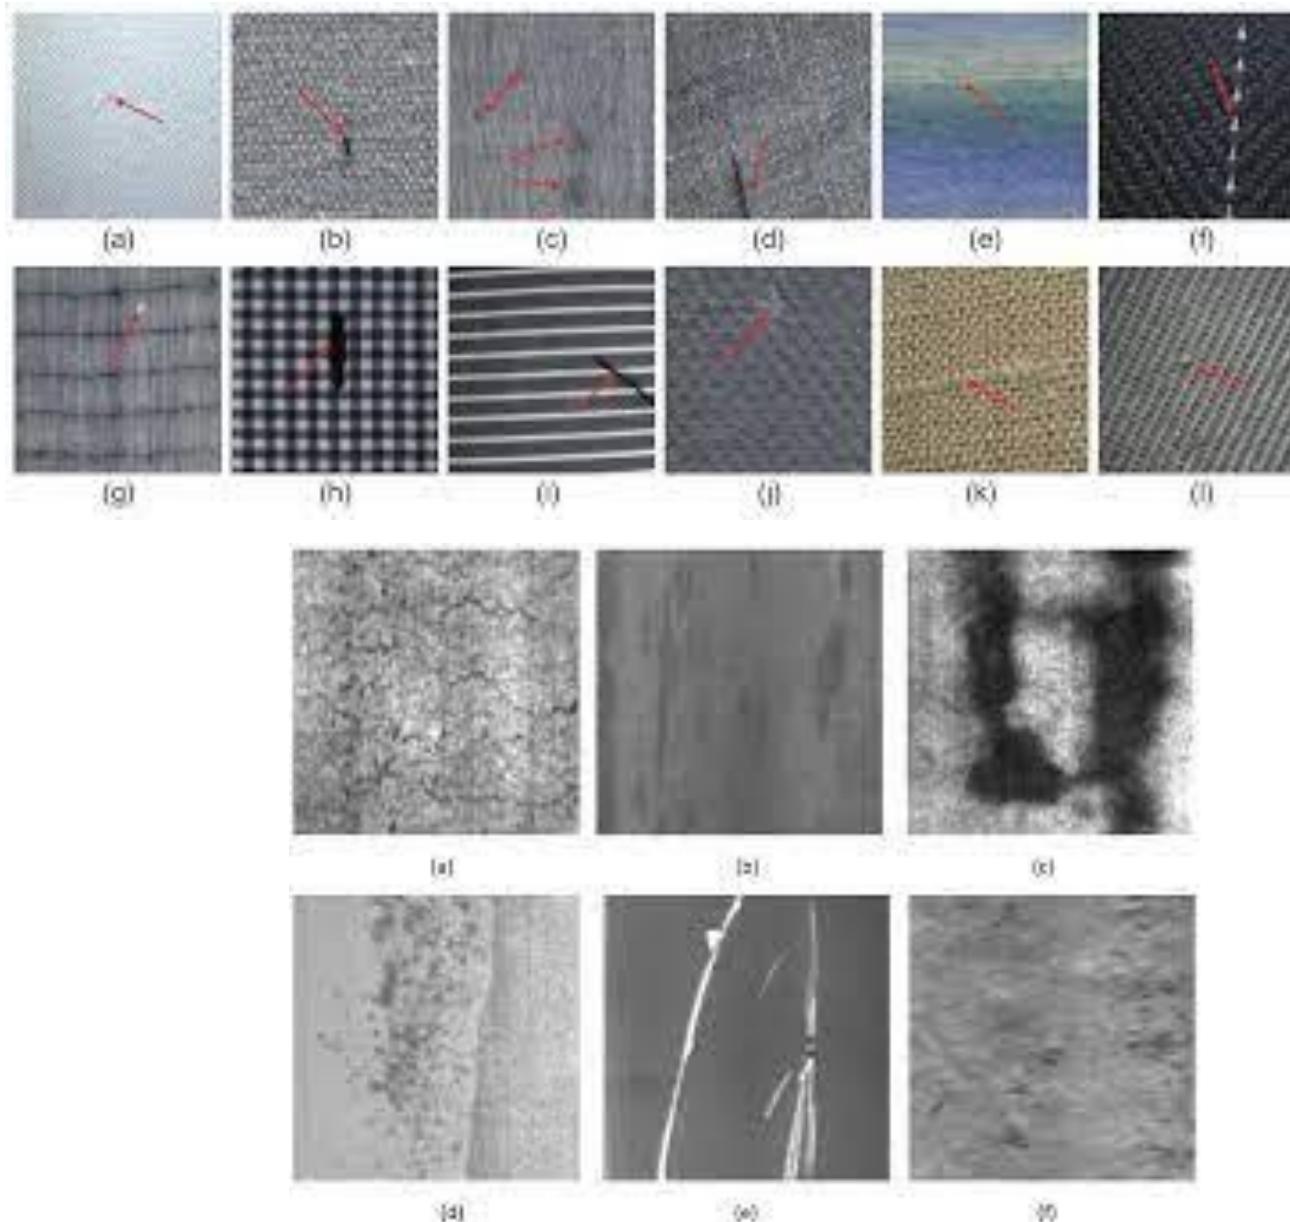


Image Annotation

Given an image, what are the words that describe the image?

Model	
Human Annotation	sky jet plane smoke
Automatic Annotation	plane jet smoke flight prop

Image Retrieval

Given a database of images and a query string (e.g. words), what are the images that are described by the words?

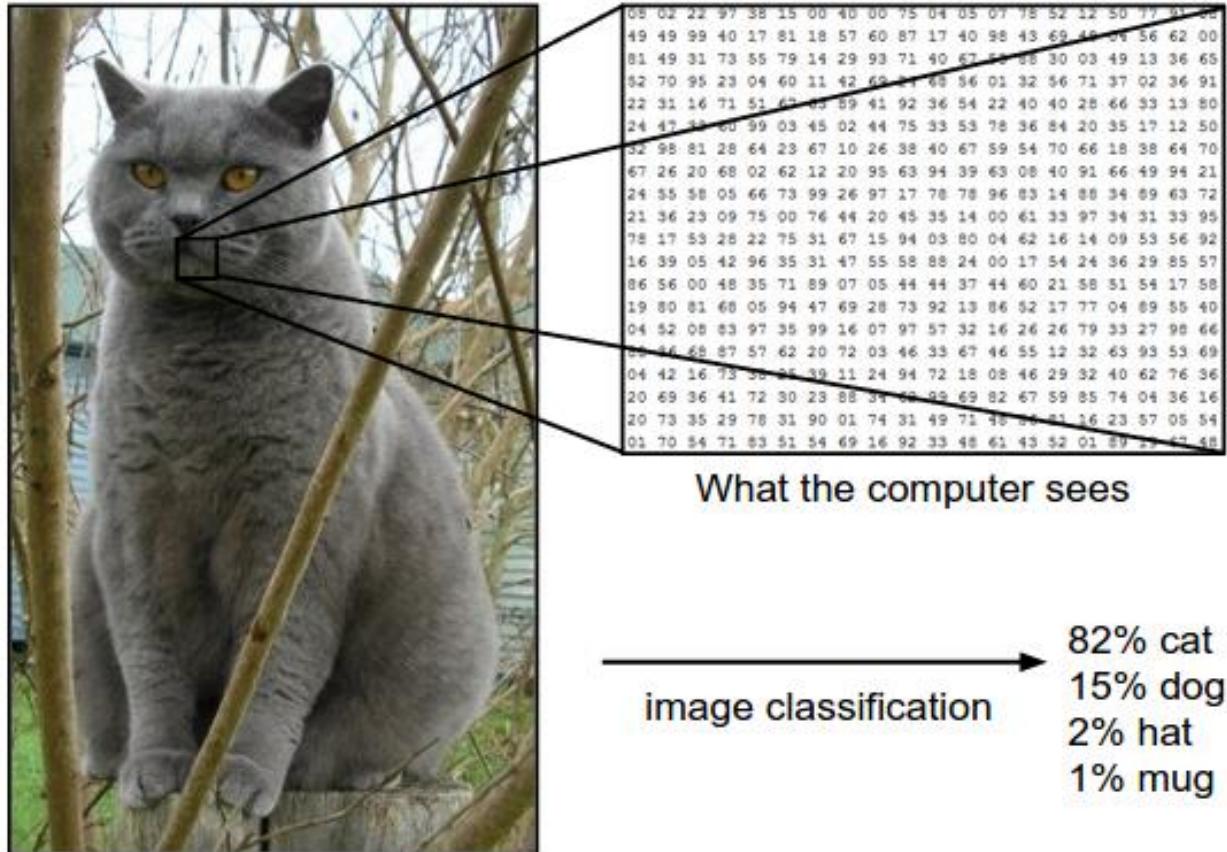
Query String: “jet”



Image classification

Image classification

Image classification is where a computer can analyze an image and identify the ‘class’ the image falls under. (Or a probability of the image being part of a ‘class’.) A class is essentially a label, for instance, ‘car’, ‘animal’, ‘building’ and so on



For example, you input an image of a cat. Image classification is the process of the computer analyzing the image and telling you it's a cat. (Or the probability that it's a cat.)

Image classification

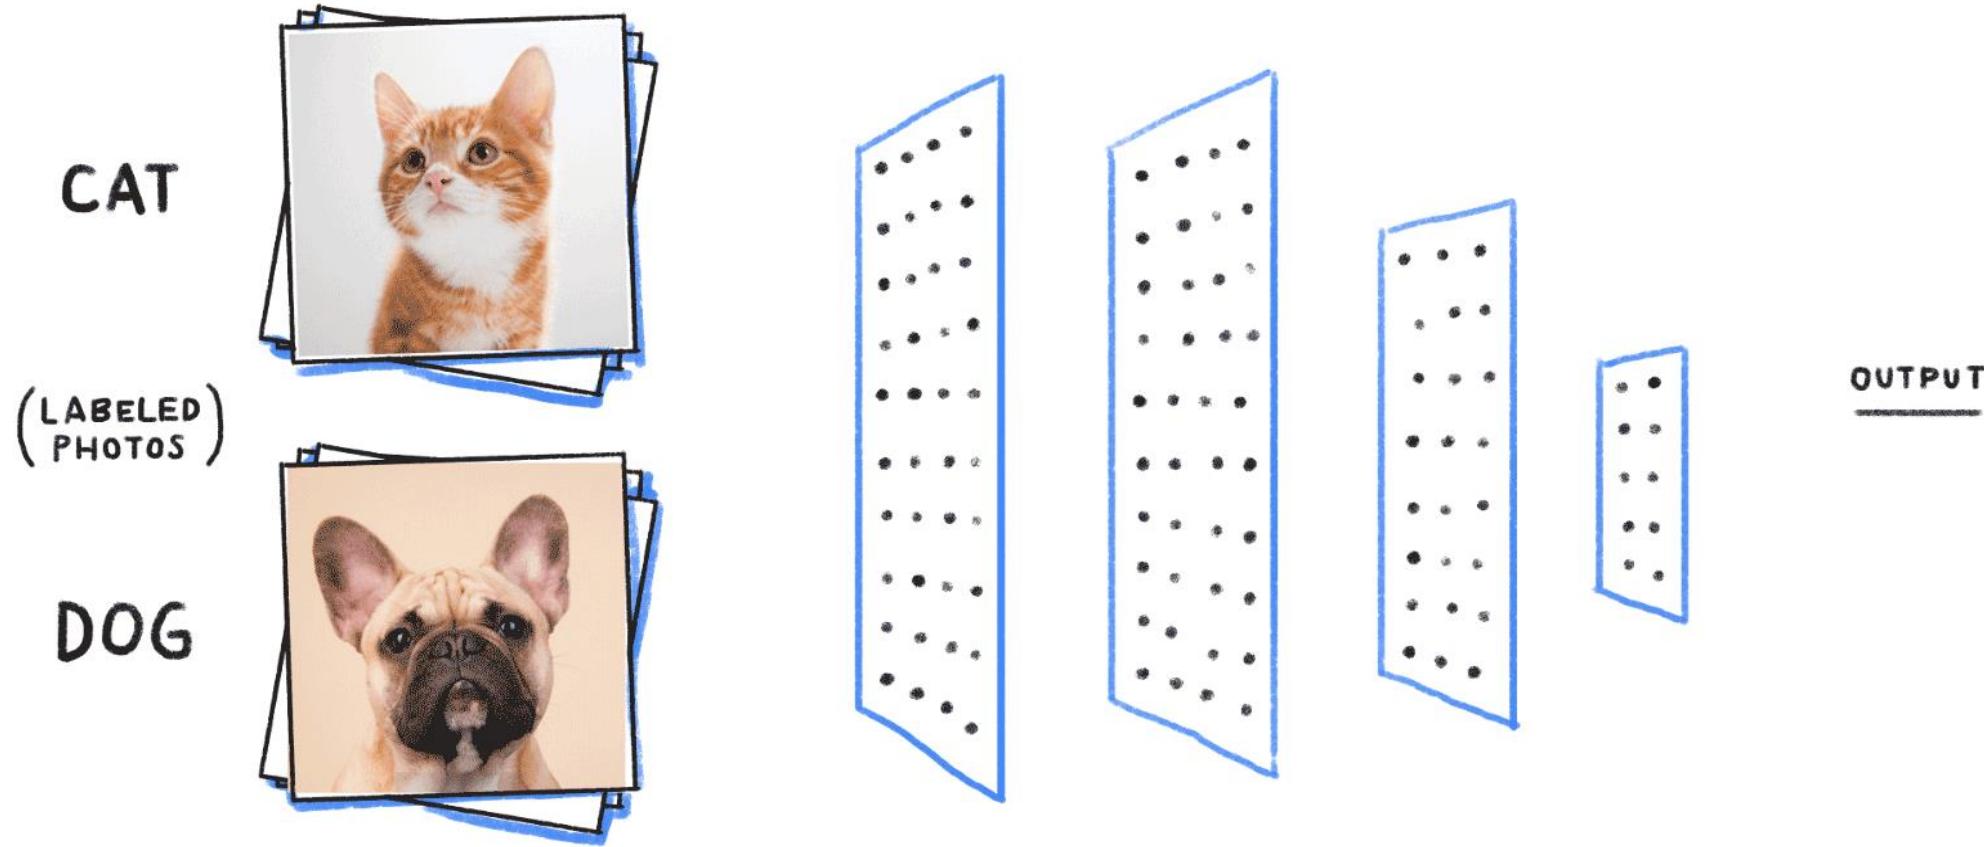
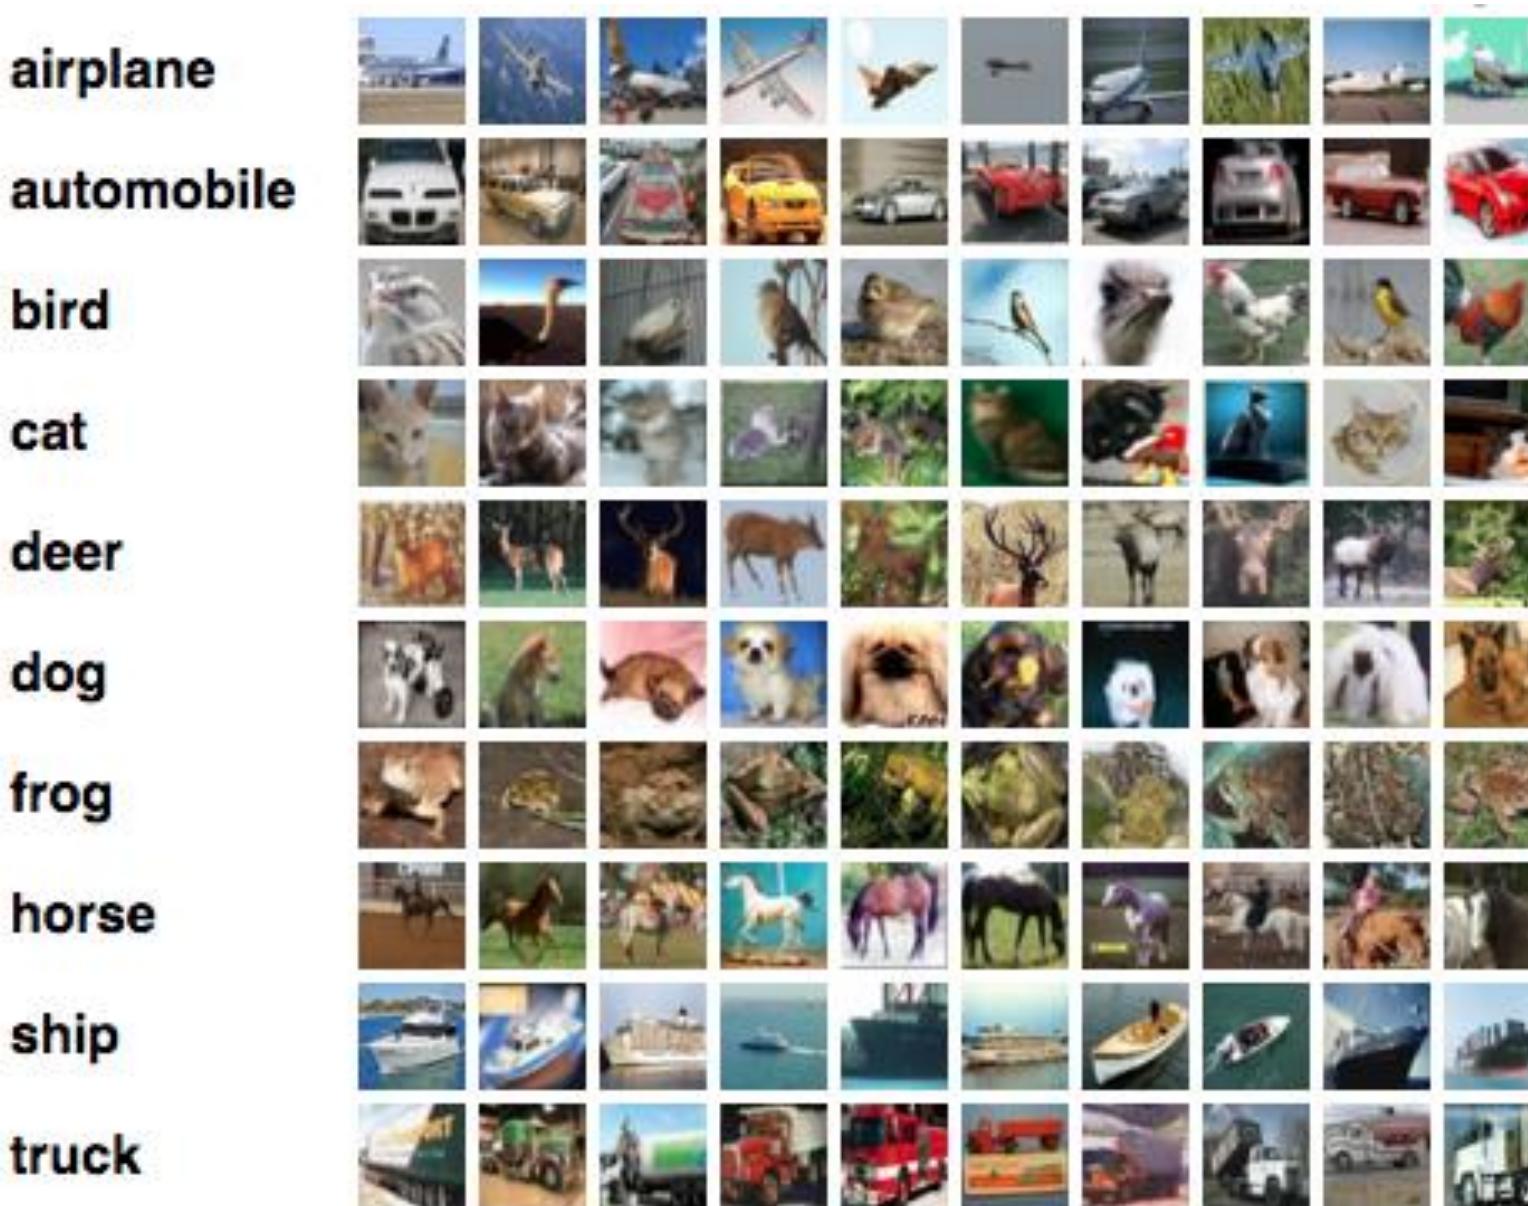


Image classification With Hand written digits

Image classification with CIFAR-10



Early Image classification

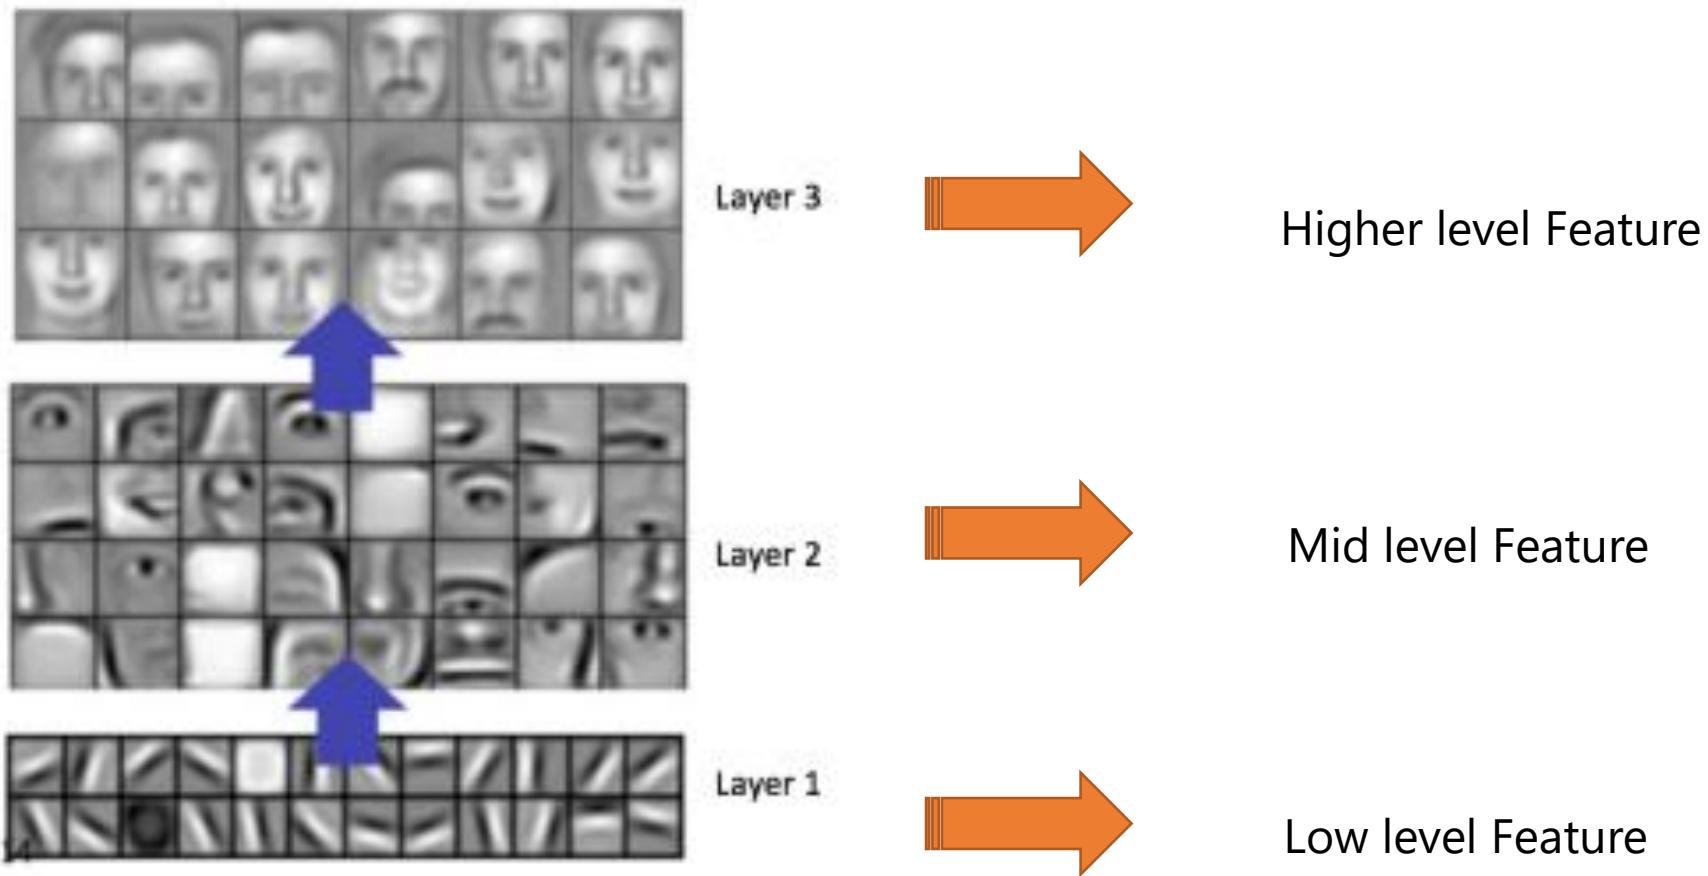
- Early image classification relied on raw pixel data.
- This meant that computers would break down images into individual pixels.
- The problem is that two pictures of the same thing can look very different.
- They can have different backgrounds, angles, poses, etcetera.
- This made it quite the challenge for computers to correctly ‘see’ and categorize images.

Adding Deep Learning for Image classification

- Deep learning involves the use of computer systems known as [neural networks](#).
- In neural networks, the input filters through hidden layers of nodes. These nodes each process the input and communicate their results to the next layer of nodes. This repeats until it reaches an output layer, and the machine provides its answer.
- Image classification with deep learning most often involves [convolutional neural networks](#), or CNNs. In CNNs, the nodes in the hidden layers don't always share their output with every node in the next layer (known as convolutional layers).

Adding Deep Learning for Image classification

- Deep learning allows machines to identify and extract features from images. This means they can learn the features to look for in images by analysing lots of pictures. So, programmers don't need to enter these filters by hand.

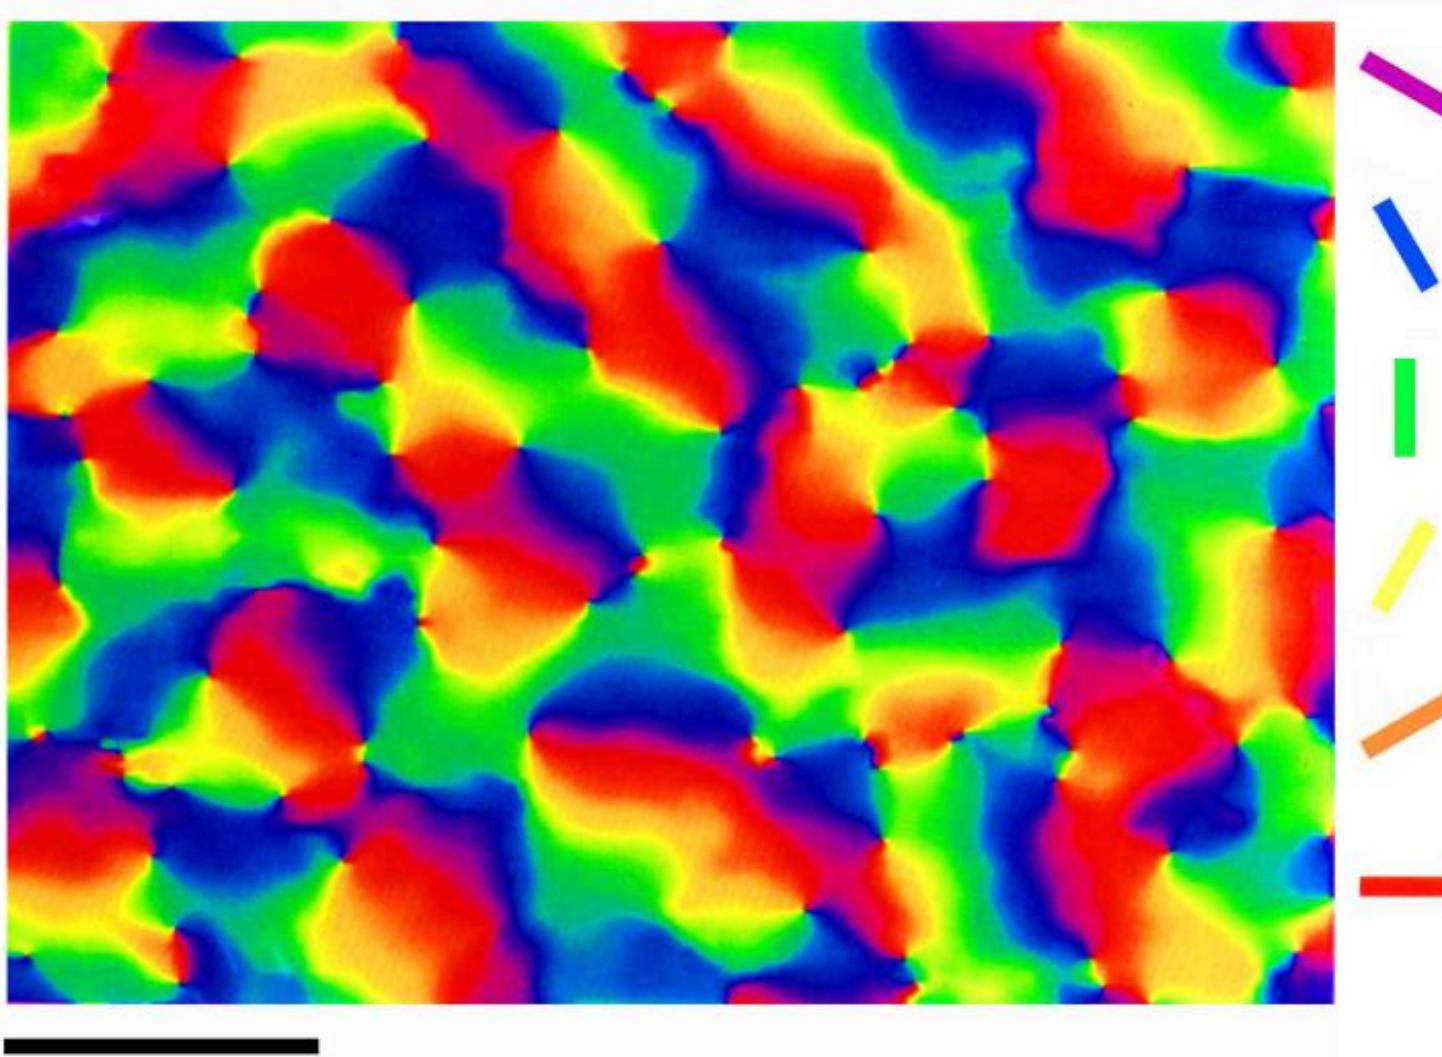


Convolutional Neural Networks

Using correlations in images

- Natural images contain spatial correlations
- For example, pixels along a contour or edge
- How can we use these correlations?

Biological inspiration



What is a convolution?

```
array = np.array([0, 0, 0, 0, 0, 1, 1, 1, 1, 1])
kernel = np.array([-1, 1])
conv = np.array([0, 0, 0, 0, 0, 0, 0, 0, 0])
conv[0] = (kernel * array[0:2]).sum()
conv[1] = (kernel * array[1:3]).sum()
conv[2] = (kernel * array[2:4]).sum()
...
for ii in range(8):
    conv[ii] = (kernel * array[ii:ii+2]).sum()
conv
```

```
array([0, 0, 0, 0, 1, 0, 0, 0, 0])
```

Convolution in one dimension

```
array = np.array([0, 0, 1, 1, 0, 0, 1, 1, 0, 0])
kernel = np.array([-1, 1])
conv = np.array([0, 0, 0, 0, 0, 0, 0, 0, 0])
for ii in range(8):
    conv[ii] = (kernel * array[ii:ii+2]).sum()
conv
```

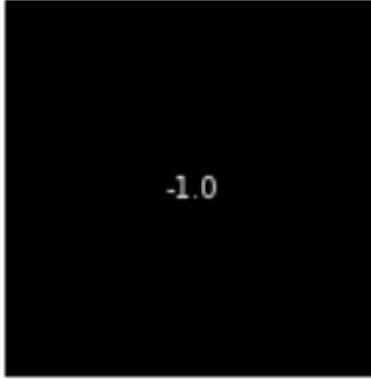
```
array([ 0,  1,  0, -1,  0,  1,  0, -1,  0])
```

Image convolution



-1.0

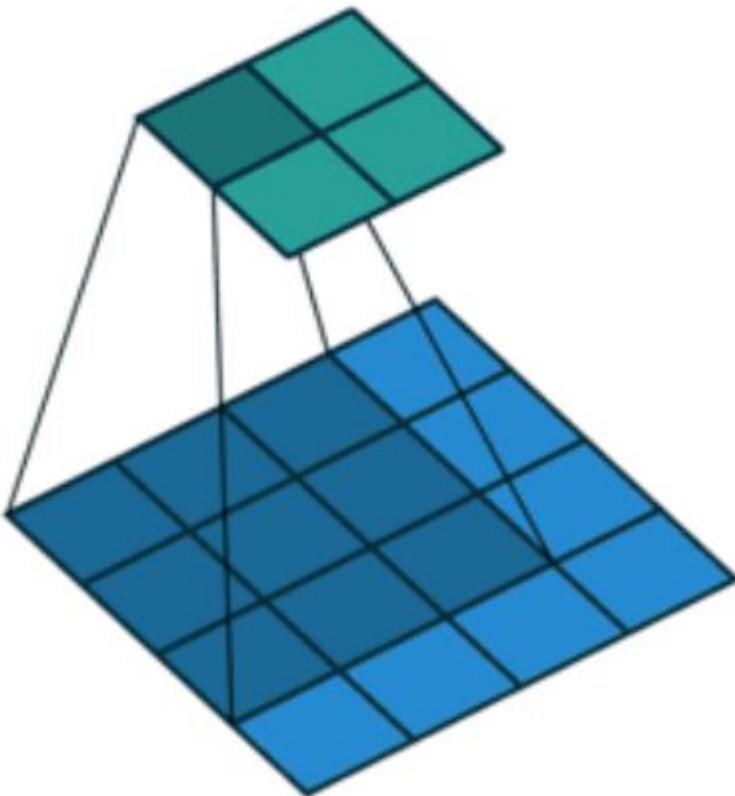
10



-1.0

Image convolution

Convolution



1	1	1	0	0
0	1	1	1	0
0	0	1 _{x1}	1 _{x0}	1 _{x1}
0	0	1 _{x0}	1 _{x1}	0 _{x0}
0	1	1 _{x1}	0 _{x0}	0 _{x1}

Image

4	3	4
2	4	3
2	3	4

Convolved
Feature

Convoluting a $5 \times 5 \times 1$ image with a $3 \times 3 \times 1$ kernel to get a $3 \times 3 \times 1$ convolved feature

Implementing convolutions in Keras

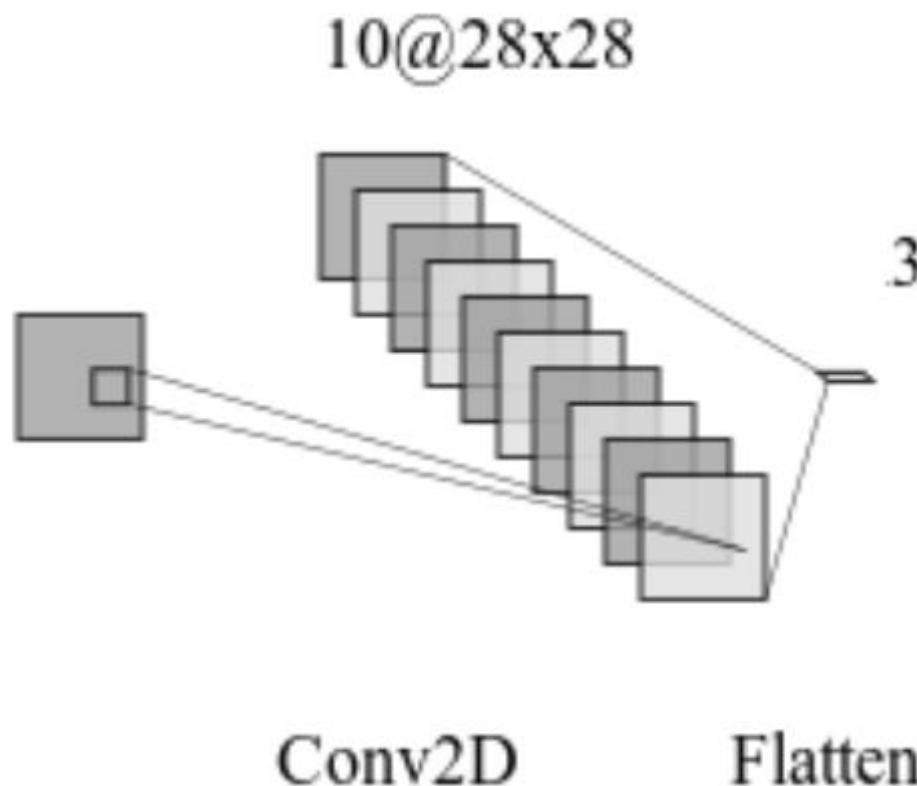
Keras Convolution layer

```
from keras.layers import Conv2D  
Conv2D(10, kernel_size=3, activation='relu')
```

Integrating convolution layers into a network

```
from keras.models import Sequential
from keras.layers import Dense, Conv2D, Flatten
model = Sequential()
model.add(Conv2D(10, kernel_size=3, activation='relu',
                input_shape=(img_rows, img_cols, 1)))
model.add(Flatten())
model.add(Dense(3, activation='softmax'))
```

Our CNN



Fitting a CNN

```
model.compile(optimizer='adam',  
              loss='categorical_crossentropy',  
              metrics=['accuracy'])  
  
train_data.shape
```

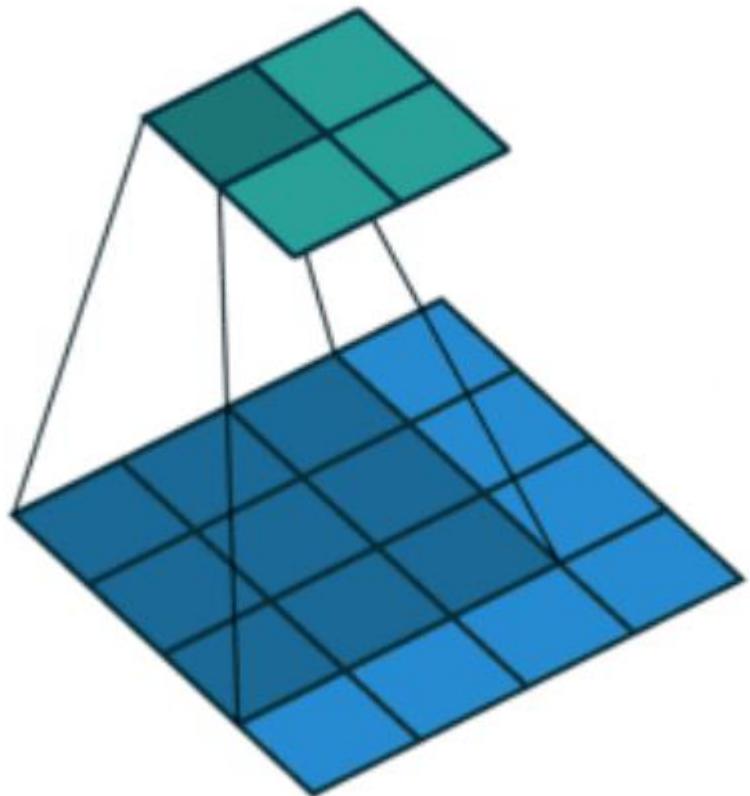
```
(50, 28, 28, 1)
```

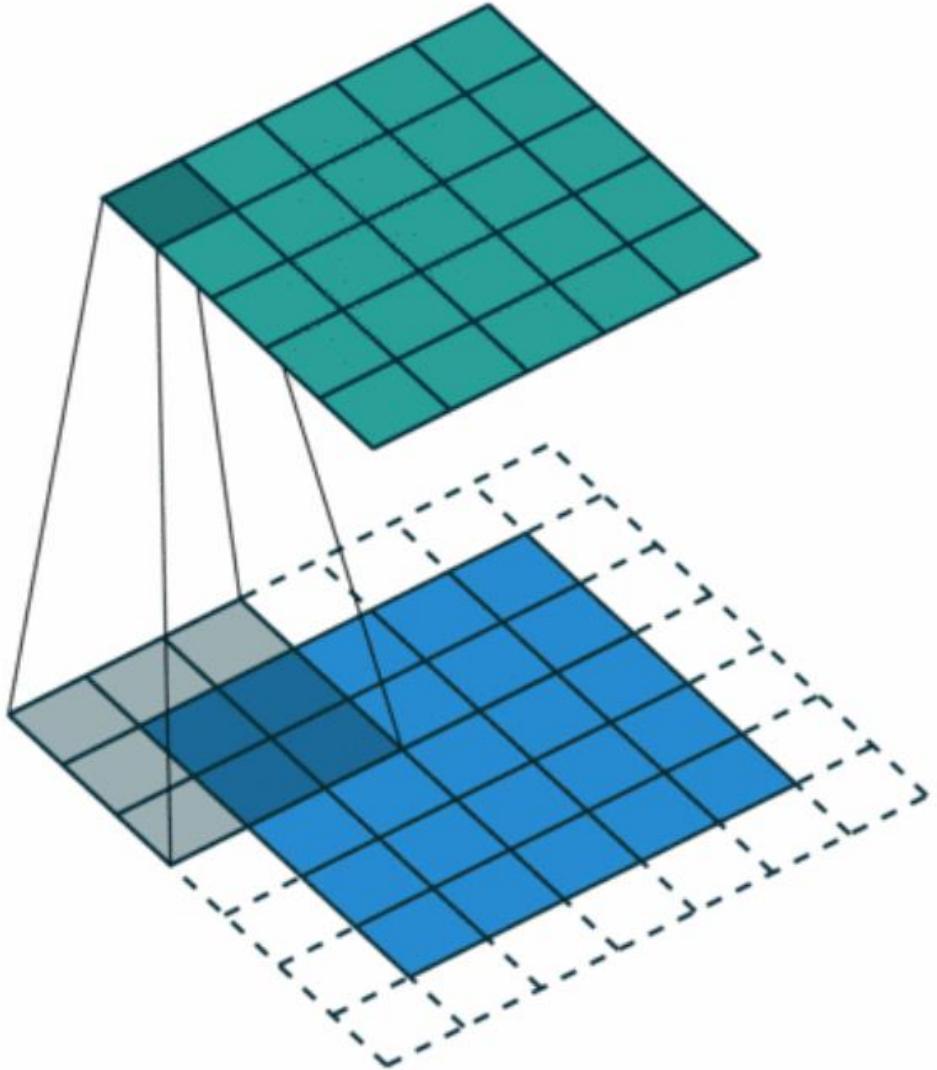
```
model.fit(train_data, train_labels, validation_split=0.2,  
          epochs=3)
```

```
model.evaluate(test_data, test_labels, epochs=3)
```

Tweaking your convolutions

Convolution





SAME padding: $5 \times 5 \times 1$ image is padded with 0s to create a $6 \times 6 \times 1$ image

Zero padding in Keras

```
model.add(Conv2D(10, kernel_size=3, activation='relu',
                 input_shape=(img_rows, img_cols, 1)),
           padding='valid')
```

Zero padding in Keras

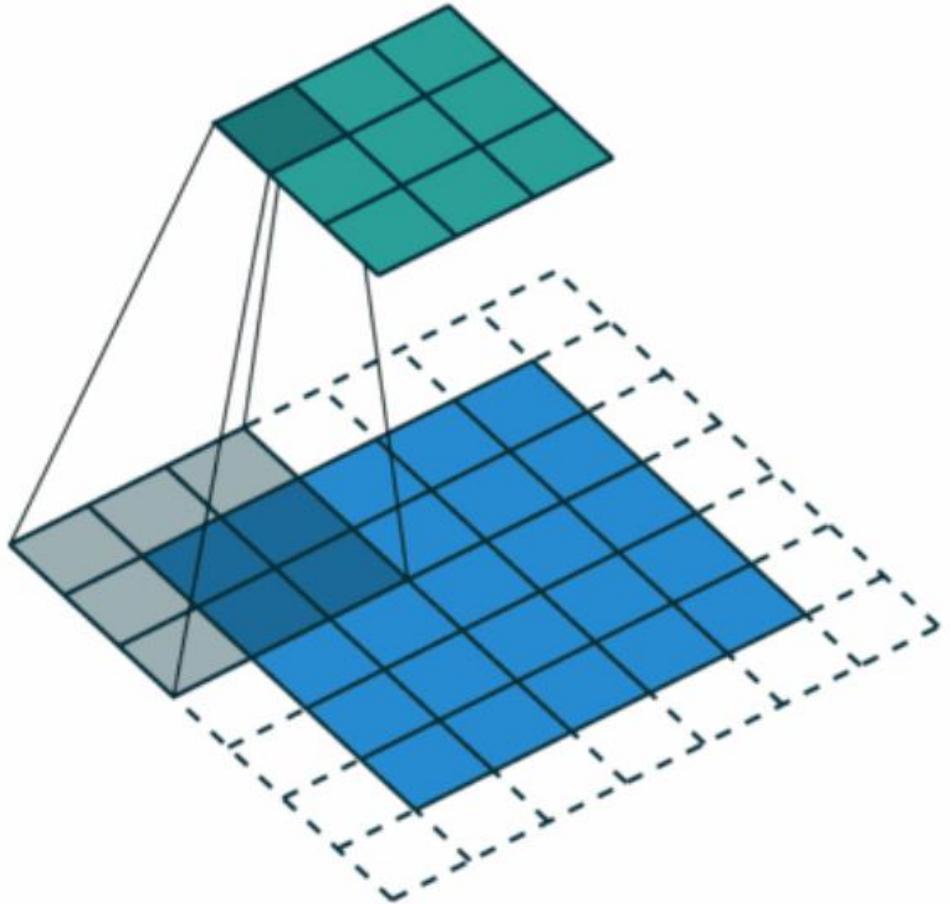
```
model.add(Conv2D(10, kernel_size=3, activation='relu',  
                 input_shape=(img_rows, img_cols, 1)),  
                 padding='same')
```

Strides in Keras

```
model.add(Conv2D(10, kernel_size=3, activation='relu',
                 input_shape=(img_rows, img_cols, 1)),
           strides=1)
```

Strides in Keras

```
model.add(Conv2D(10, kernel_size=3, activation='relu',
                 input_shape=(img_rows, img_cols, 1)),
           strides=2)
```



Convolution Operation with Stride Length = 2

Calculating the size of the output

$$O = ((I - K + 2P)/S) + 1$$

where

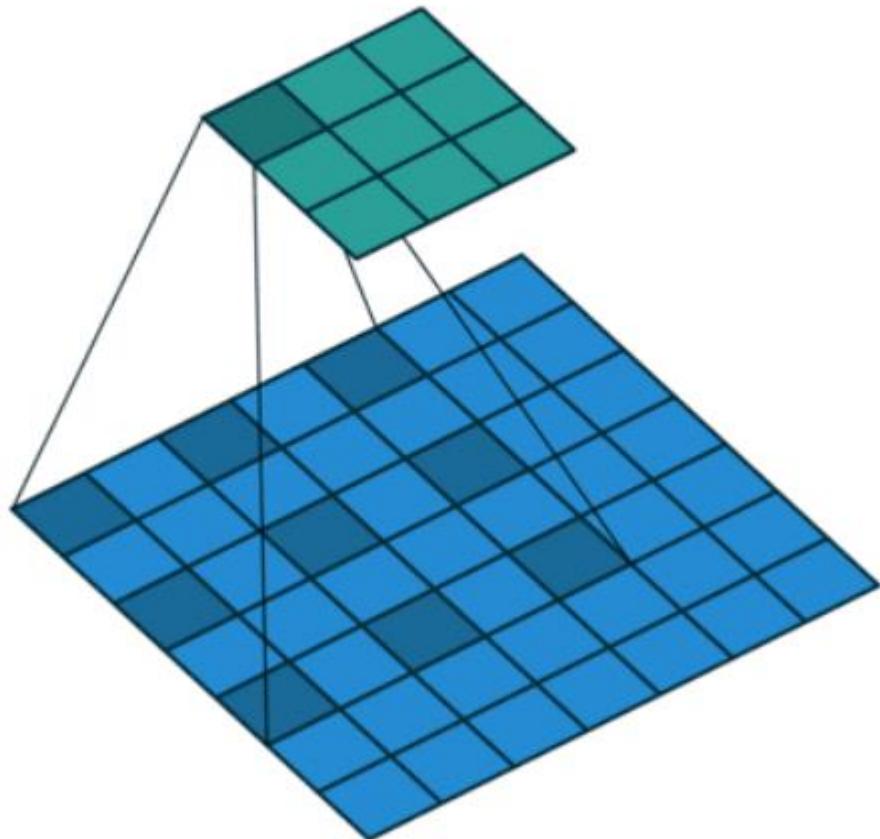
- I = size of the input
- K = size of the kernel
- P = size of the zero padding
- S = strides

Calculating the size of the output

$$28 = ((28 - 3 + 2)/1) + 1$$

$$10 = ((28 - 3 + 2)/3) + 1$$

Dilated convolutions

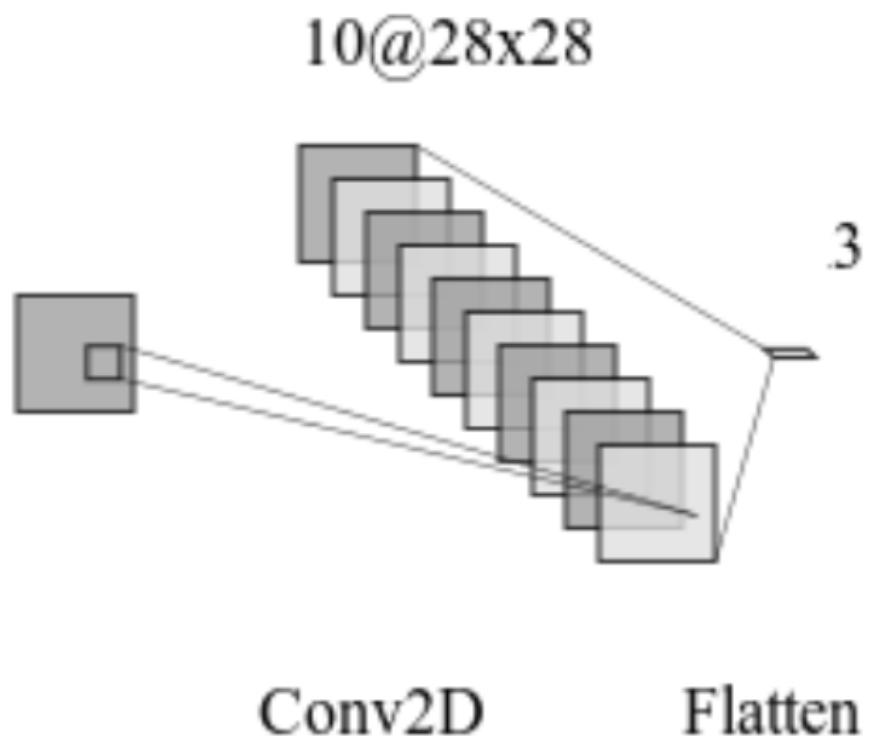


Dilation in Keras

```
model.add(Conv2D(10, kernel_size=3, activation='relu',
                 input_shape=(img_rows, img_cols, 1)),
           dilation_rate=2)
```

Going deeper

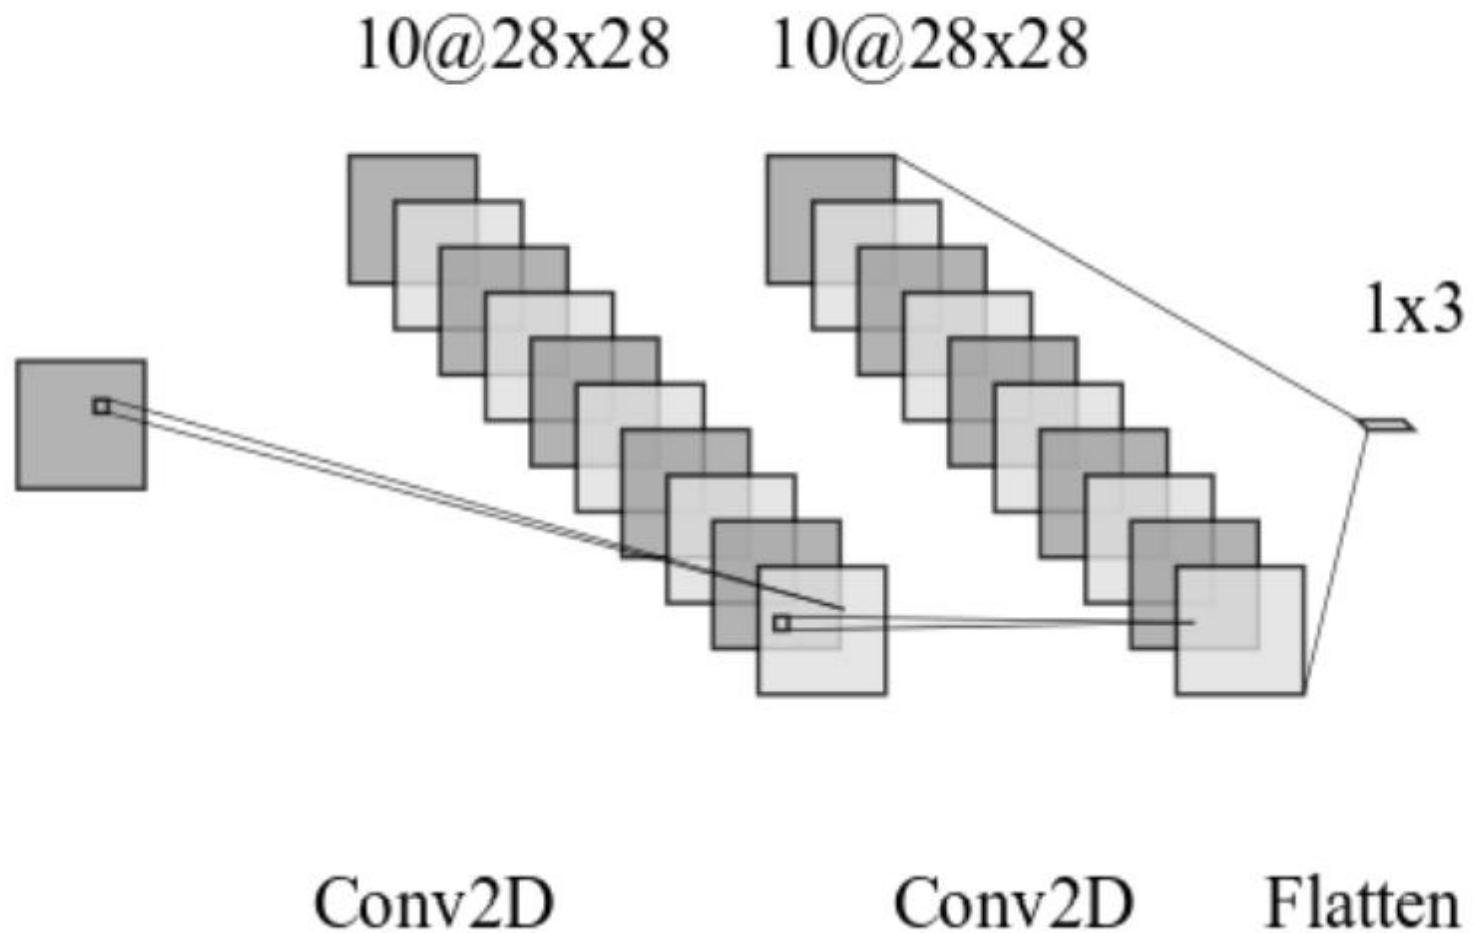
Network with one convolutional layer



Network with one convolutional layer: implementation

```
model = Sequential()  
model.add(Conv2D(10, kernel_size=2, activation='relu',  
                input_shape=(img_rows, img_cols, 1)))  
model.add(Flatten())  
model.add(Dense(3, activation='softmax'))
```

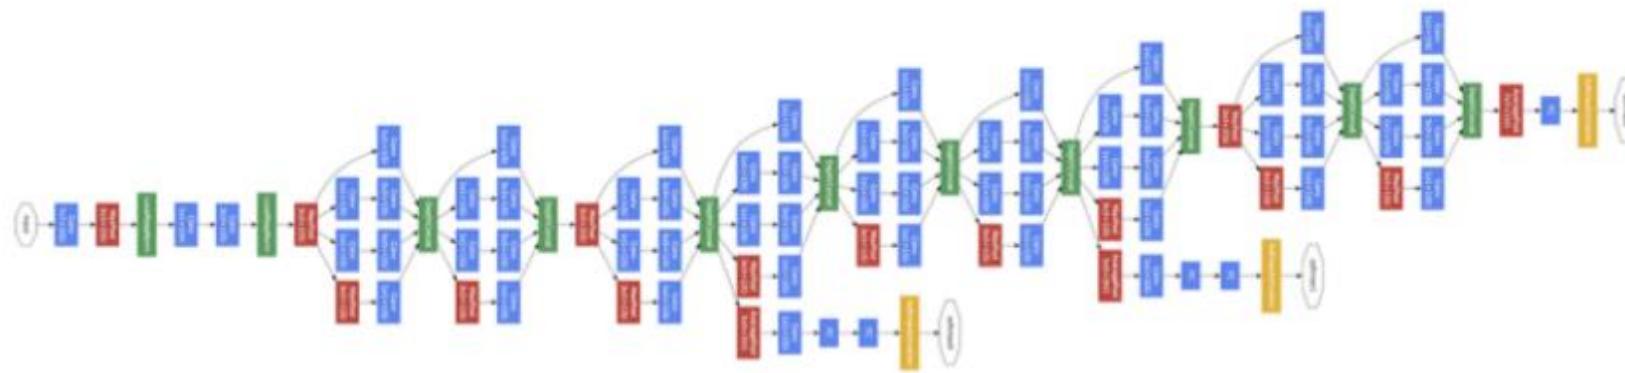
Building a deeper network



Building a deep network

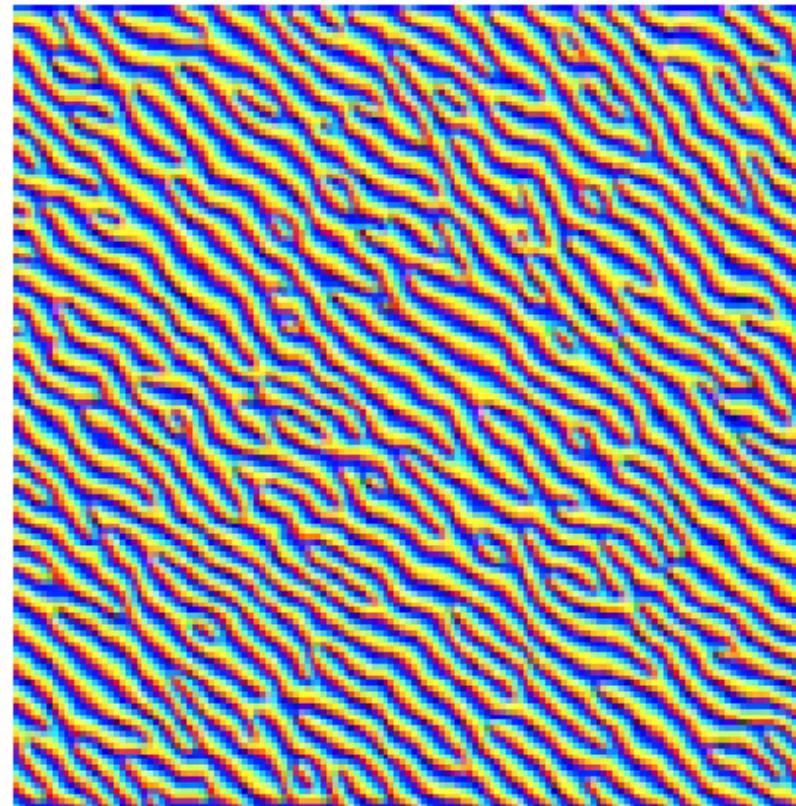
```
model = Sequential()
model.add(Conv2D(10, kernel_size=2, activation='relu',
                input_shape=(img_rows, img_cols, 1),
                padding='equal'))
# Second convolutional layer
model.add(Conv2D(10, kernel_size=2, activation='relu'))
model.add(Flatten())
model.add(Dense(3, activation='softmax'))
```

Why do we want deep networks?

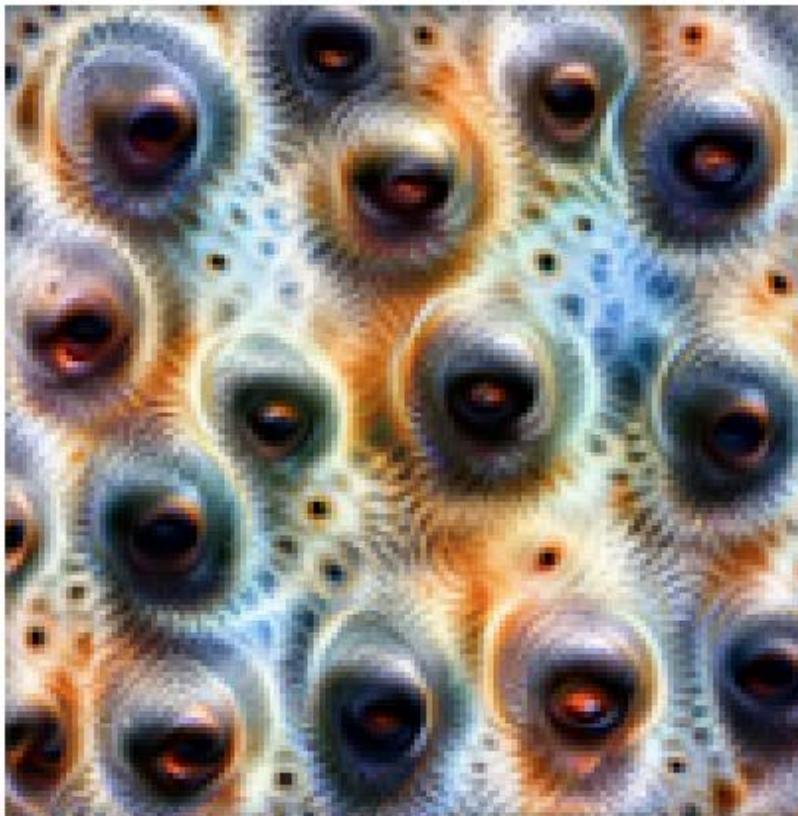
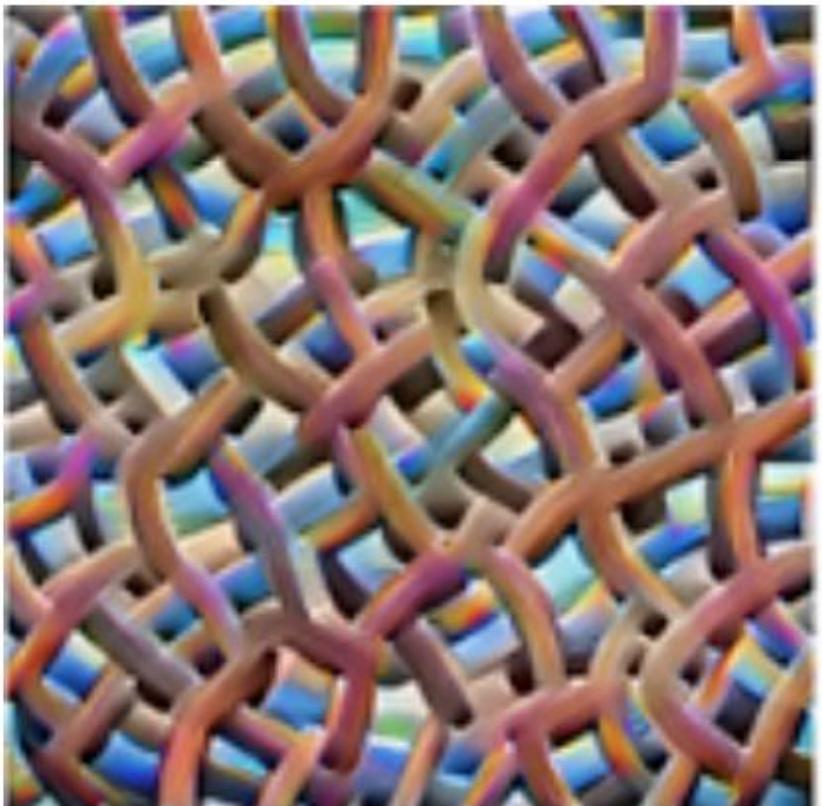


Convolution
Pooling
Softmax
Other

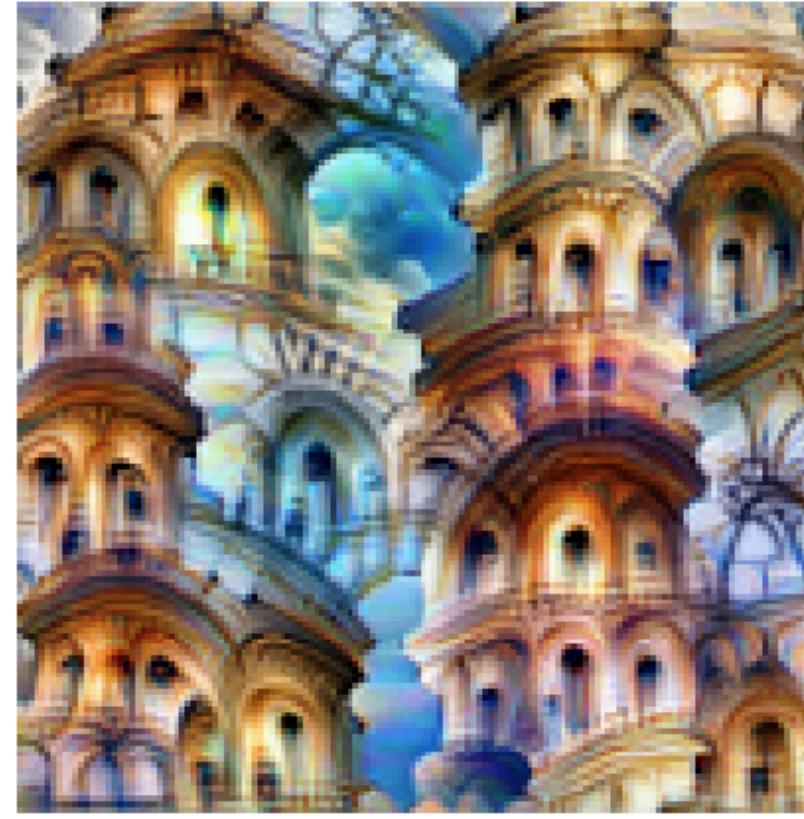
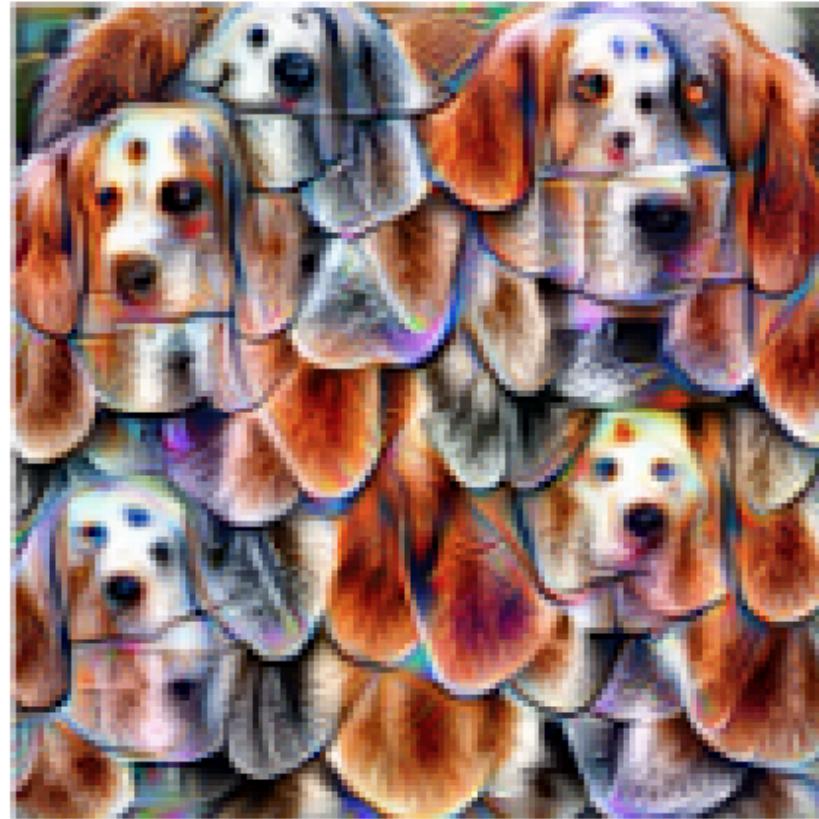
Features in early layers

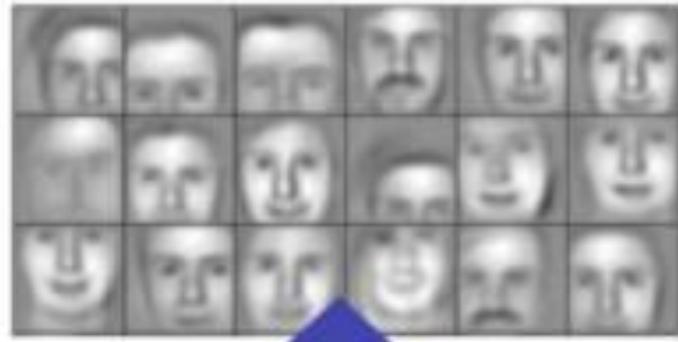


Features in intermediate layers



Features in late layers

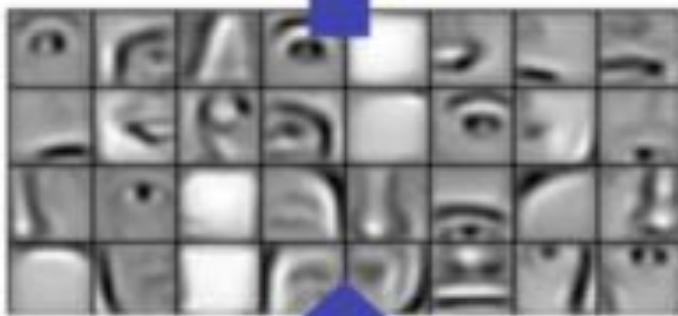




Layer 3



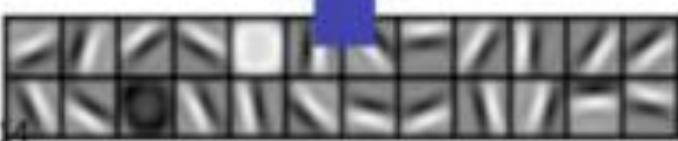
Higher level Feature



Layer 2



Mid level Feature



Layer 1



Low level Feature

How deep?

- Depth comes at a computational cost
- May require more data

How many parameters?

Counting parameters

```
model = Sequential()  
  
model.add(Dense(10, activation='relu',  
               input_shape=(784,)))  
  
model.add(Dense(10, activation='relu'))  
  
model.add(Dense(3, activation='softmax'))
```

```
# Call the summary method  
model.summary()
```

Layer (type)	Output Shape	Param #
dense_1 (Dense)	(None, 10)	7850
dense_2 (Dense)	(None, 10)	110
dense_3 (Dense)	(None, 3)	33
Total params:	7,993	
Trainable params:	7,993	
Non-trainable params:	0	

Counting parameters

```
model.add(Dense(  
    10, activation='relu',  
    input_shape=(784,)))
```

$$\begin{aligned} \text{parameters} &= 784 * 10 + 10 \\ &= 7850 \end{aligned}$$

```
model.add(Dense(  
    10, activation='relu'))
```

$$\begin{aligned} \text{parameters} &= 10 * 10 + 10 \\ &= 110 \end{aligned}$$

```
model.add(Dense(  
    3, activation='softmax'))
```

$$\begin{aligned} \text{parameters} &= 10 * 3 + 3 \\ &= 33 \end{aligned}$$

$$7850 + 110 + 33 = 7993$$

```
model.summary()
```

Layer (type)	Output Shape	Param #
<hr/>		
dense_1 (Dense)	(None, 10)	7850

dense_2 (Dense)	(None, 10)	110
-----------------	------------	-----

dense_3 (Dense)	(None, 3)	33
<hr/>		

Total params: 7,993

Trainable params: 7,993

Non-trainable params: 0

The number of parameters in a CNN

```
model = Sequential()

model.add(Conv2D(10, kernel_size=3, activation='relu',
                input_shape=(28, 28, 1), padding='same'))

model.add(Conv2D(10, kernel_size=3, activation='relu',
                padding='same'))

model.add(Flatten())

model.add(Dense(3, activation='softmax'))
```

```
model.summary()
```

Layer (type)	Output Shape	Param #
<hr/>		
conv2d_1 (Conv2D)	(None, 28, 28, 10)	100
<hr/>		
conv2d_2 (Conv2D)	(None, 28, 28, 10)	910
<hr/>		
flatten_3 (Flatten)	(None, 7840)	0
<hr/>		
dense_4 (Dense)	(None, 3)	23523
<hr/>		
Total params: 24,533		
Trainable params: 24,533		
Non-trainable params: 0		
<hr/>		

The number of parameters in a CNN

```
model.add(  
    Conv2D(10, kernel_size=3,  
           activation='relu',  
           input_shape=(28, 28, 1),  
           padding='same'))
```

```
model.add(  
    Conv2D(10, kernel_size=3,  
           activation='relu',  
           padding='same'))
```

```
model.add(Flatten())
```

```
model.add(Dense(  
    3, activation='softmax'))
```

$$\text{parameters} = 9 * 10 + 10$$

$$= 100$$

$$\text{parameters} = 10 * 9 * 10 + 10$$

$$= 910$$

$$\text{parameters} = 0$$

$$\text{parameters} = 7840 * 3 + 3$$

$$= 23523$$

$$100 + 910 + 0 + 23523 = 24533$$

Increasing the number of units in each layer

```
model = Sequential()

model.add(Dense(5, activation='relu',
               input_shape=(784,), padding='same'))

model.add(Dense(15, activation='relu', padding='same'))

model.add(Dense(3, activation='softmax'))
```

```
model.summary()
```

Layer (type)	Output Shape	Param #
=====		
dense_1 (Dense)	(None, 5)	3925
=====		
dense_2 (Dense)	(None, 15)	90
=====		
dense_3 (Dense)	(None, 3)	48
=====		

Total params: 4,063

Trainable params: 4,063

Non-trainable params: 0

Increasing the number of units in each layer

```
model = Sequential()

model.add(Conv2D(5, kernel_size=3, activation='relu',
                input_shape=(28, 28, 1),
                padding="same"))

model.add(Conv2D(15, kernel_size=3, activation='relu',
                padding="same"))

model.add(Flatten())

model.add(Dense(3, activation='softmax'))
```

```
model.summary()
```

Layer (type)	Output Shape	Param #
<hr/>		
conv2d_12 (Conv2D)	(None, 28, 28, 5)	50
<hr/>		
conv2d_13 (Conv2D)	(None, 28, 28, 15)	690
<hr/>		
flatten_6 (Flatten)	(None, 11760)	0
<hr/>		
dense_9 (Dense)	(None, 3)	35283
<hr/>		
Total params: 36,023		
Trainable params: 36,023		
Non-trainable params: 0		

Reducing parameters with pooling

```
model.summary()
```

Layer (type)	Output Shape	Param #
<hr/>		
conv2d_12 (Conv2D)	(None, 28, 28, 5)	50

conv2d_13 (Conv2D)	(None, 28, 28, 15)	690
<hr/>		

flatten_6 (Flatten)	(None, 11760)	0
<hr/>		

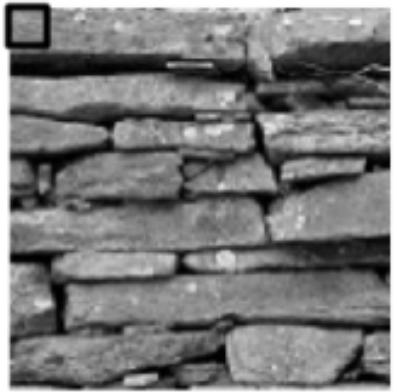
dense_9 (Dense)	(None, 3)	35283
<hr/>		

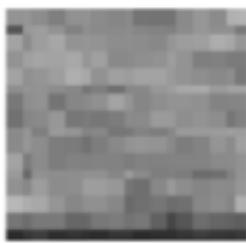
Total params: 36,023

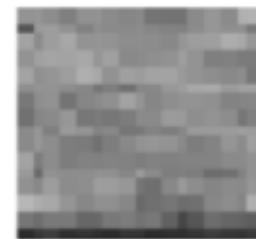
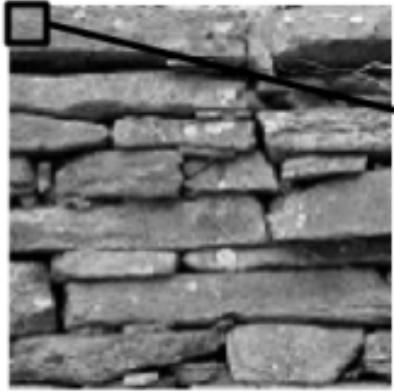
Trainable params: 36,023

Non-trainable params: 0

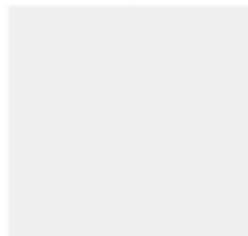


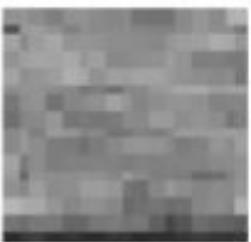
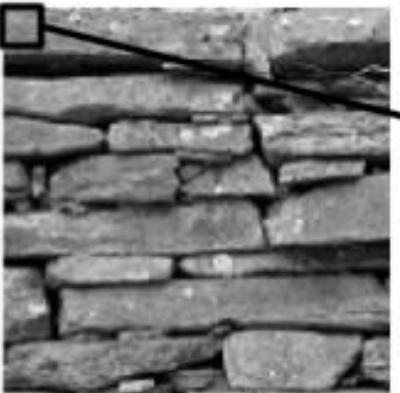




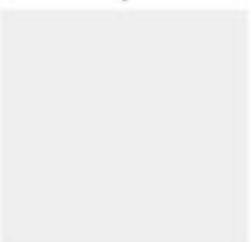


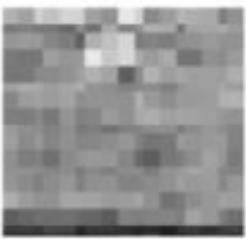
Max





Max





Max





Pooling Layer

3.0	3.0	3.0
3.0	3.0	3.0
3.0	2.0	3.0

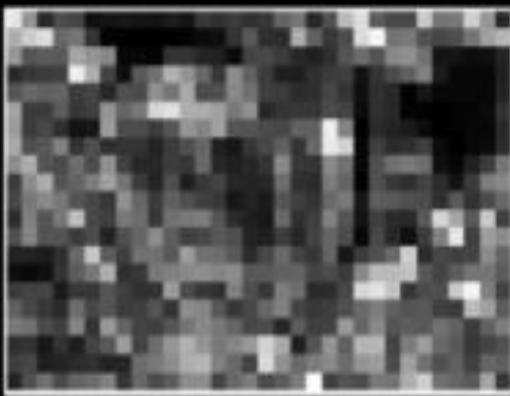
3	3	2	1	0
0	0	1	3	1
3	1	2	2	3
2	0	0	2	2
2	0	0	0	1



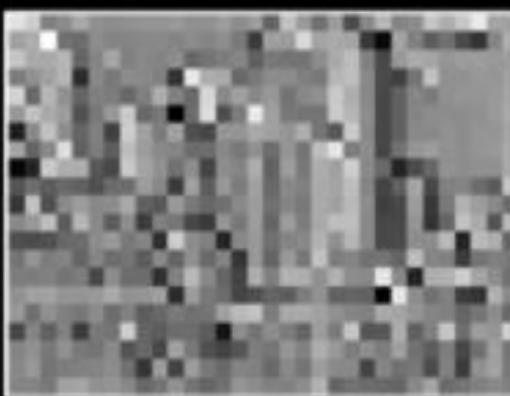
Rectified Feature Map

Pooling
→

Max



Sum

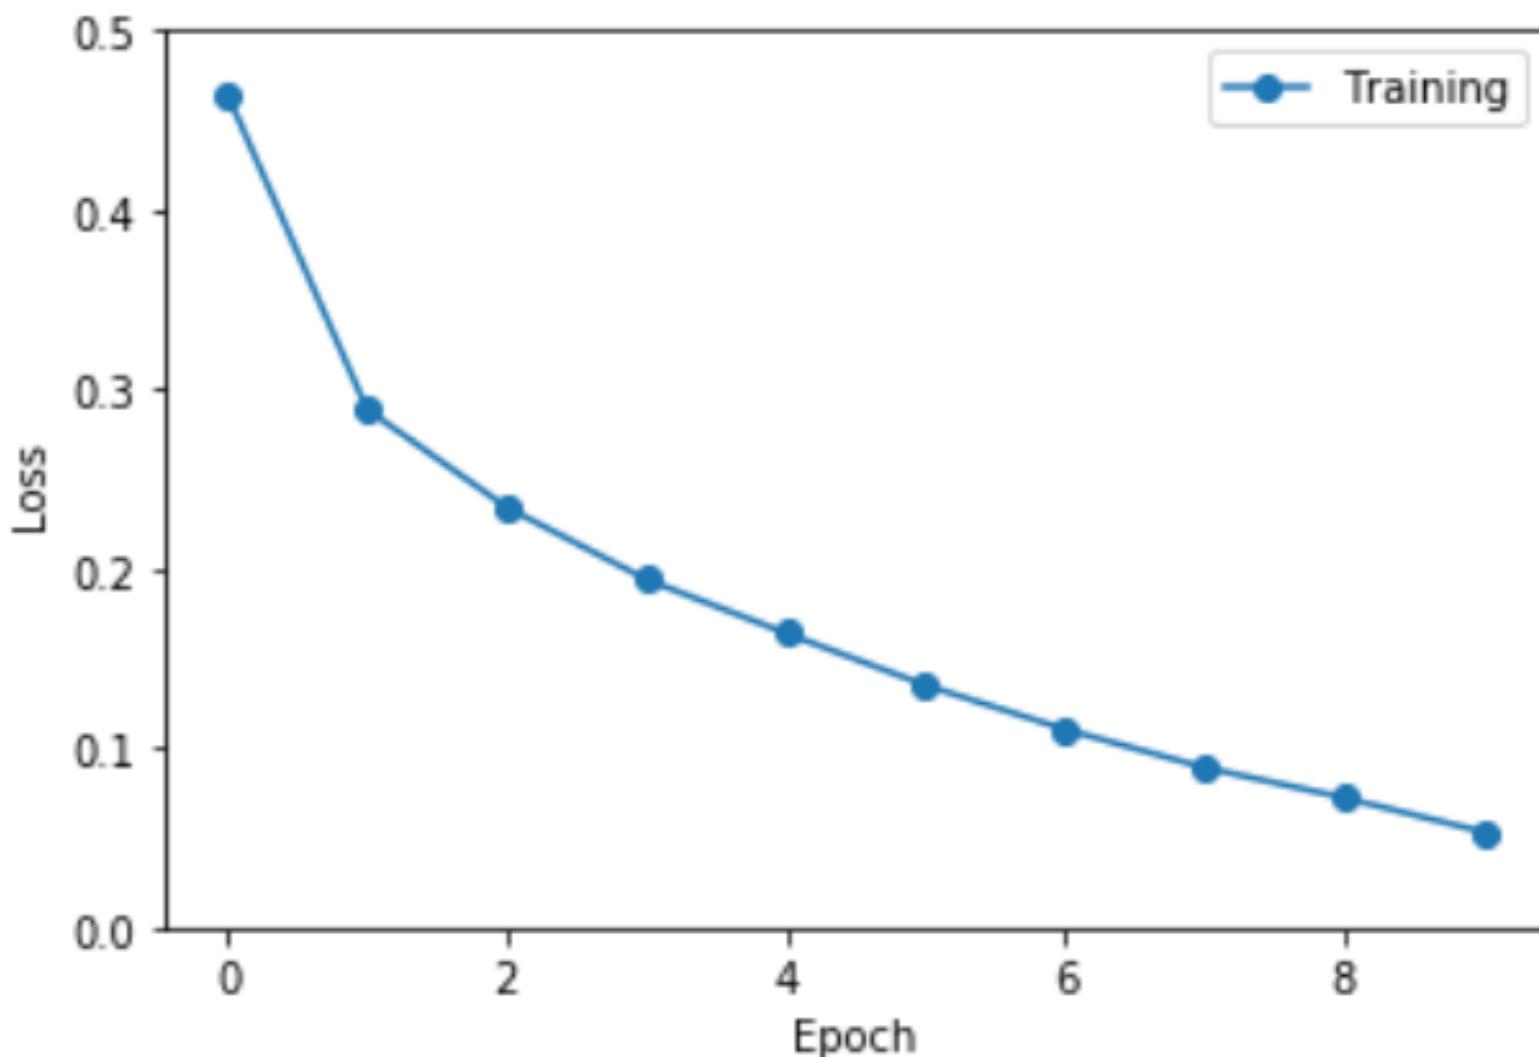


Max pooling in Keras

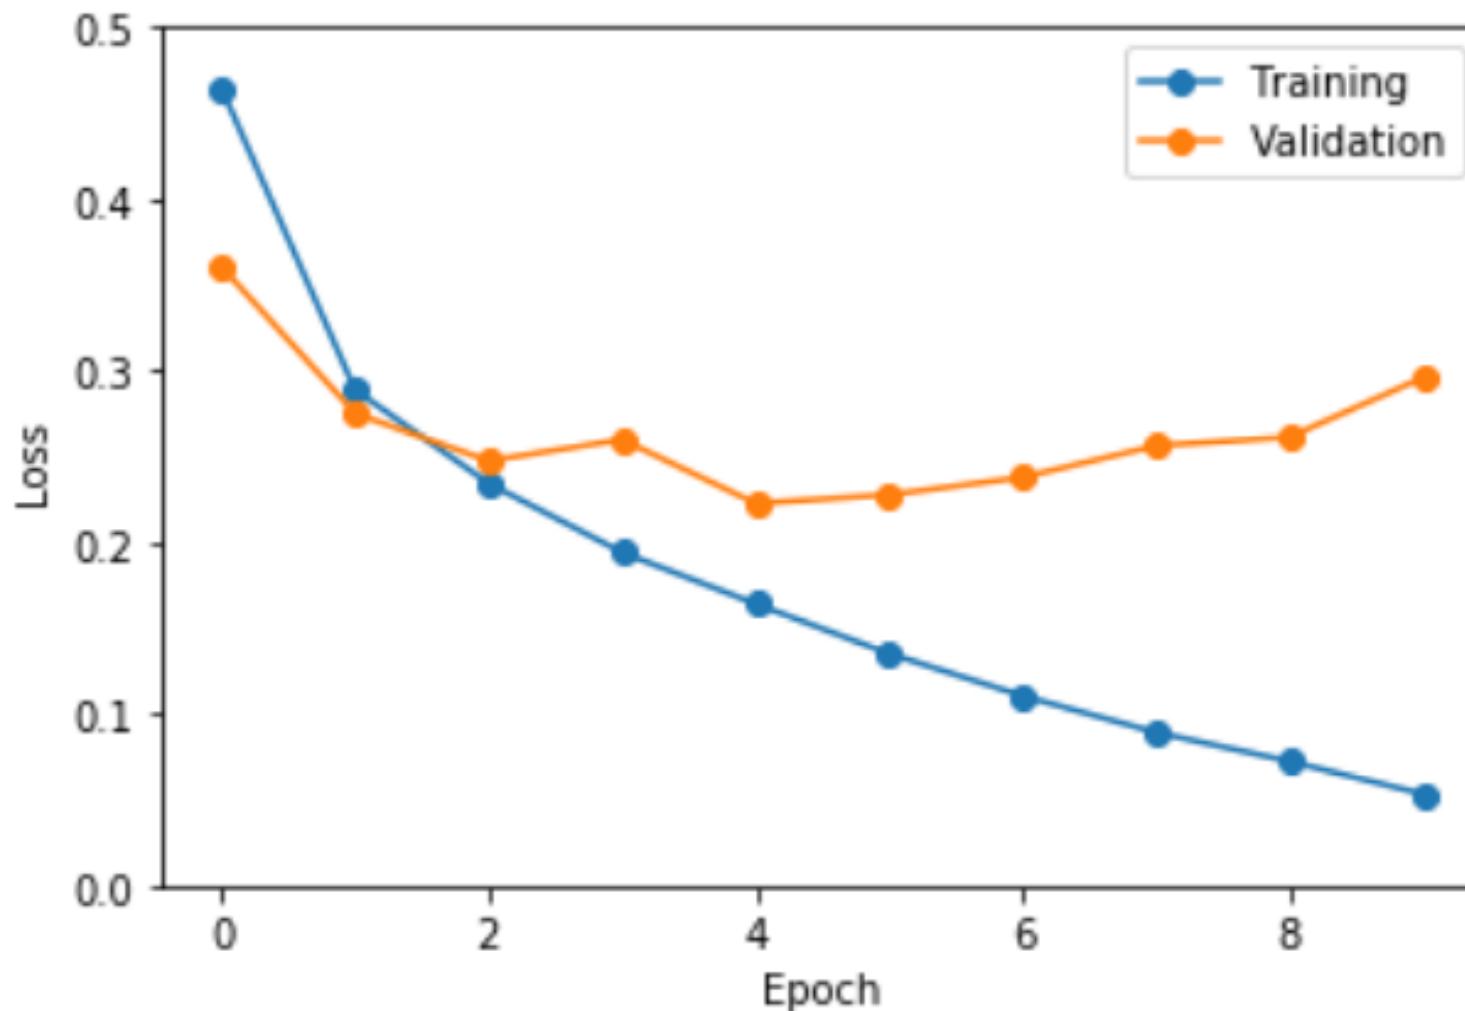
```
from keras.models import Sequential
from keras.layers import Dense, Conv2D, Flatten, MaxPool2D
model = Sequential()
model.add(Conv2D(5, kernel_size=3, activation='relu',
                input_shape=(img_rows, img_cols, 1)))
model.add(MaxPool2D(2))
model.add(Conv2D(15, kernel_size=3, activation='relu',
                input_shape=(img_rows, img_cols, 1)))
model.add(MaxPool2D(2))
model.add(Flatten())
model.add(Dense(3, activation='softmax'))
```

Tracking learning

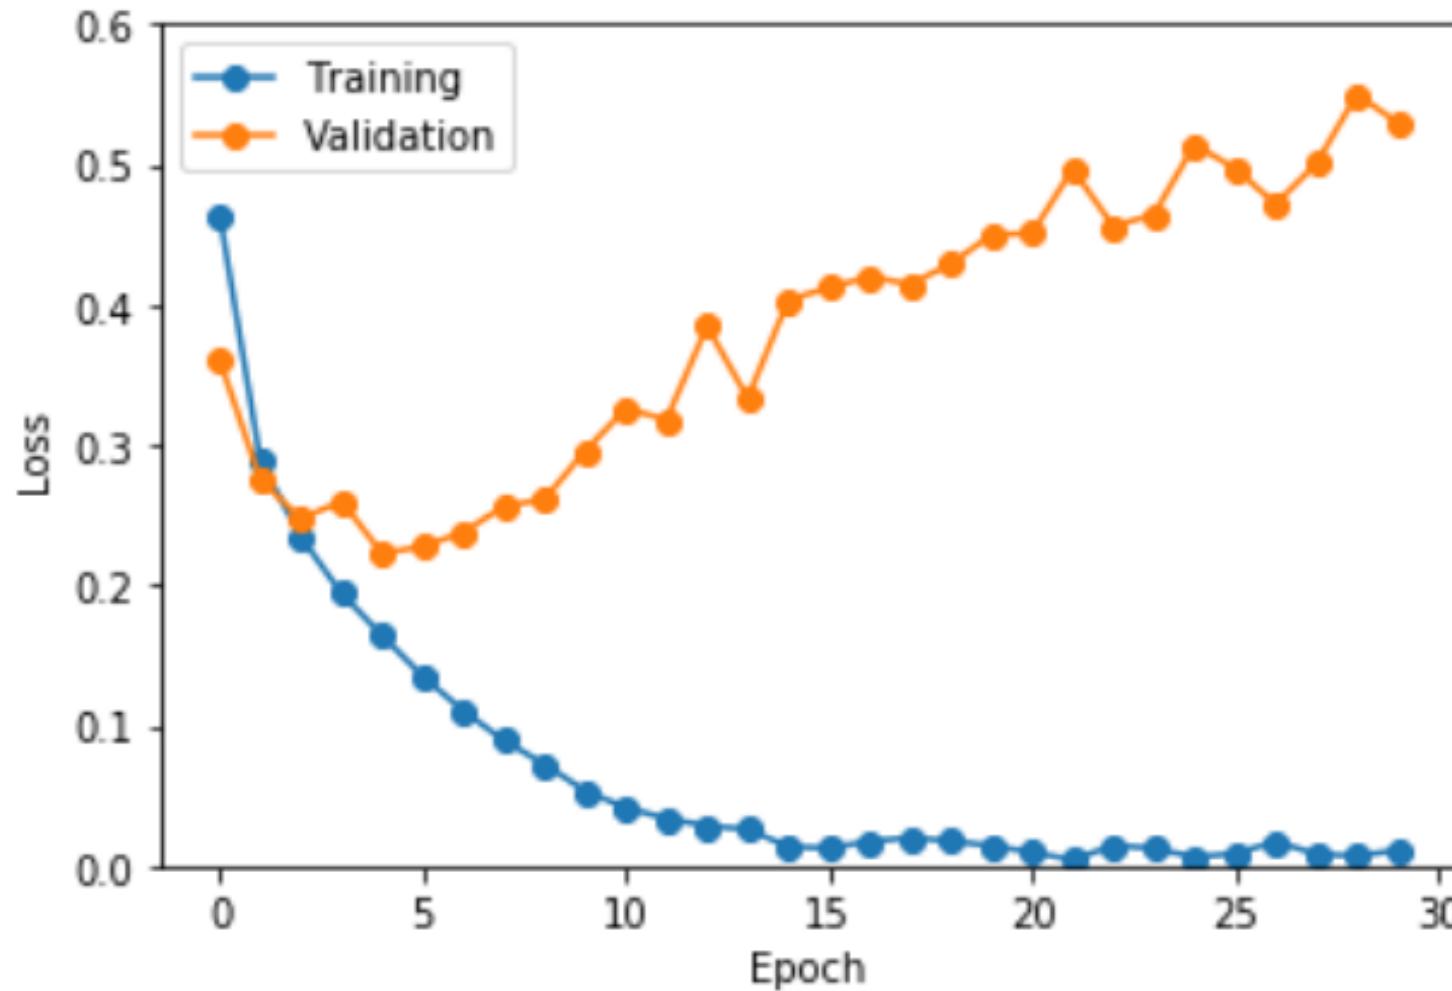
Learning curves: training



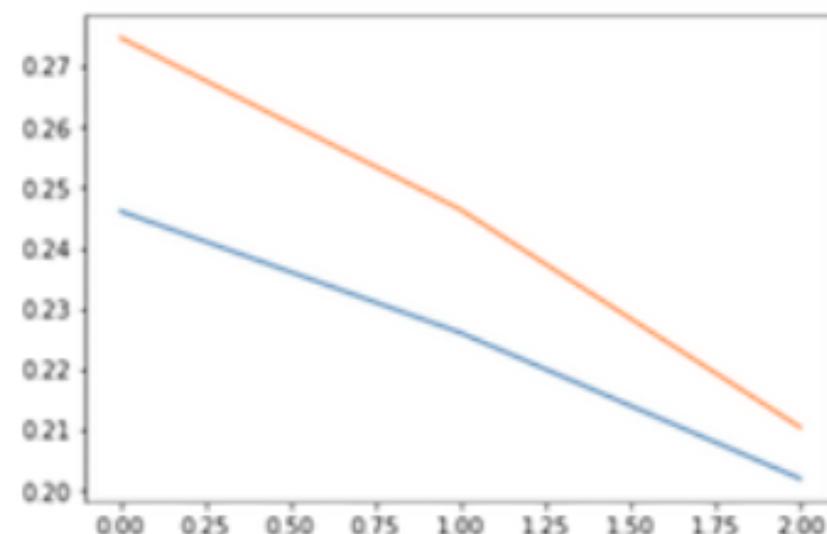
Learning curves: validation



Learning curves: overfitting



```
training = model.fit(train_data, train_labels,  
                     epochs=3, validation_split=0.2)  
  
import matplotlib.pyplot as plt  
plt.plot(training.history['loss'])  
plt.plot(training.history['val_loss'])  
plt.show()
```



Storing the optimal parameters

```
from keras.callbacks import ModelCheckpoint

# This checkpoint object will store the model parameters
# in the file "weights.hdf5"
checkpoint = ModelCheckpoint('weights.hdf5', monitor='val_loss',
                             save_best_only=True)
# Store in a list to be used during training
callbacks_list = [checkpoint]
# Fit the model on a training set, using the checkpoint as a
#callback
model.fit(train_data, train_labels, validation_split=0.2,
          epochs=3, callbacks=callbacks_list)
```

Loading stored parameters

```
model.load_weights('weights.hdf5')
model.predict_classes(test_data)
array([2, 2, 1, 2, 0, 1, 0, 1, 2, 0])
```

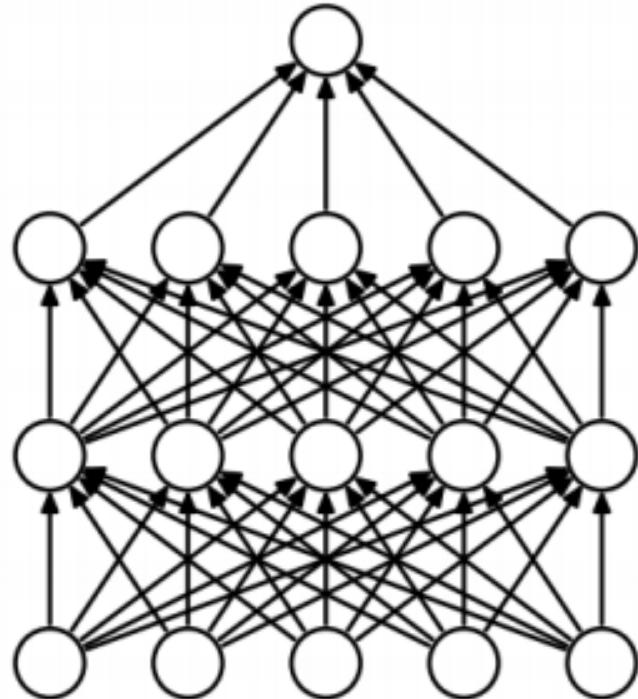
Neural network regularization

Dropout

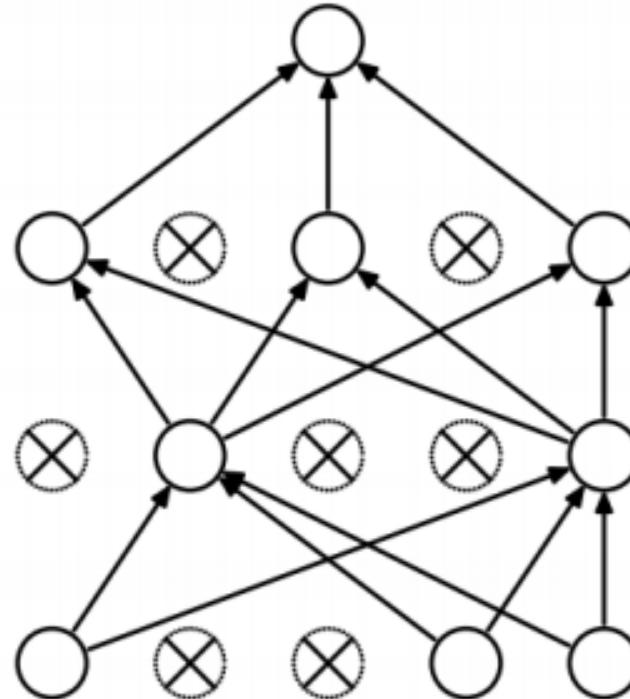
In each learning step:

- Select a subset of the units
- Ignore it in the forward pass
- And in the back-propagation of error

Dropout



(a) Standard Neural Net



(b) After applying dropout.

Dropout in Keras

```
from keras.models import Sequential
from keras.layers import Dense, Conv2D, Flatten, Dropout
model = Sequential()
model.add(Conv2D(5, kernel_size=3, activation='relu',
                input_shape=(img_rows, img_cols, 1)))
model.add(Dropout(0.25))
model.add(Conv2D(15, kernel_size=3, activation='relu'))
model.add(Flatten())
model.add(Dense(3, activation='softmax'))
```

Batch normalization

- Rescale the outputs

Batch Normalization in Keras

```
from keras.models import Sequential
from keras.layers import Dense, Conv2D, Flatten, BatchNormalization

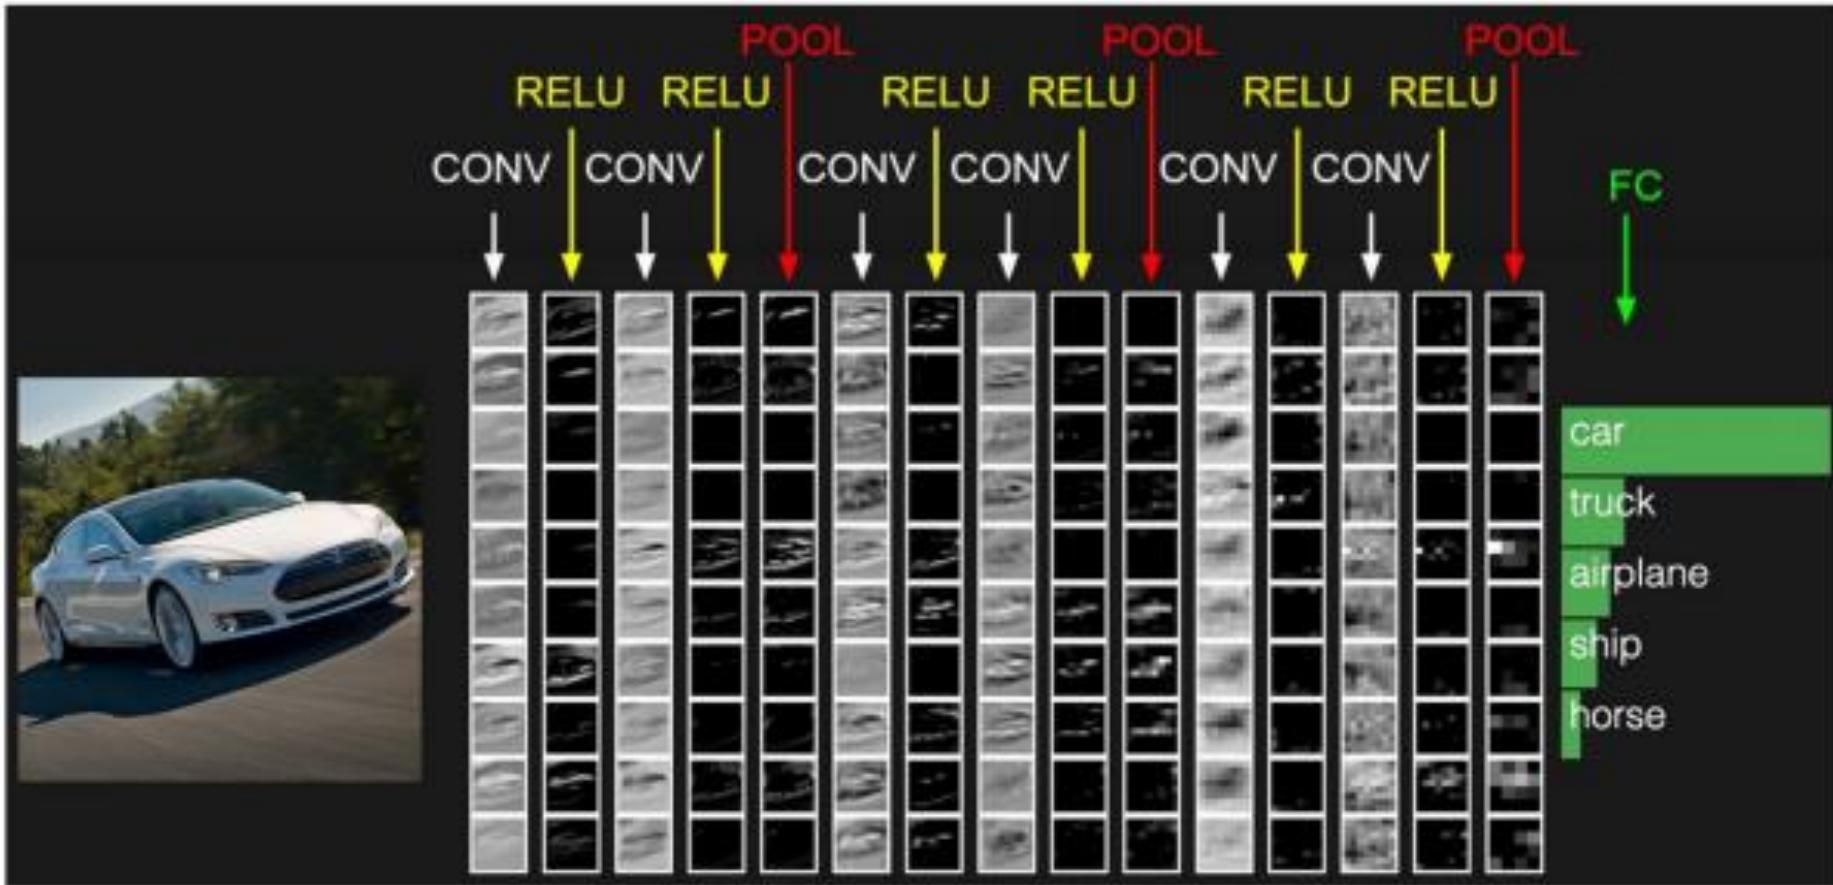
model = Sequential()
model.add(Conv2D(5, kernel_size=3, activation='relu',
                input_shape=(img_rows, img_cols, 1)))
model.add(BatchNormalization())
model.add(Conv2D(15, kernel_size=3, activation='relu'))
model.add(Flatten())
model.add(Dense(3, activation='softmax'))
```

Be careful when using them together!

The disharmony between dropout and batch normalization

Interpreting the model

CNN Visualization



What did we learn?

- Image classification
- Convolutions
- Reducing the number of parameters
 - Tweaking your convolutions
 - Adding pooling layers
- Improving your network
 - Regularization
- Understanding your network
 - Monitoring learning
 - Interpreting the parameters